

Catamaran/ARIFIC: Architecture Research in FPGAs in the Cloud

Tutorial at HPCA-29, Montreal

**Rishiyur S. Nikhil
Bluespec, Inc.
Sunday, February 26, 2023**



What is Catamaran/ARIFIC?

Architecture Researcher's Focus

- CPU microarchitecture
- Memory systems (caches, MMUs, coherence, WMMs, ...)
- Accelerators

RISC-V
CPU

...

RISC-V
CPU

Tightly-coupled
accelerator

Loosely-
coupled
accelerator

Memory subsystem
(caches, MMUs, ...)

Catamaran/ARIFIC (infrastructure in the cloud)

Host computer running Linux

System Control:

- Load ELF/memhex (programs, data) into DDR
- Assert/Deassert Core's RESET
- Configure CPU
- Collect stats, traces

Console
(to/from UART)

Disk-device
support

Network-device
support

GDB control
of CPU on
FPGA

PCIe
connection

FPGA

UART
(for console)

Disk-device
support

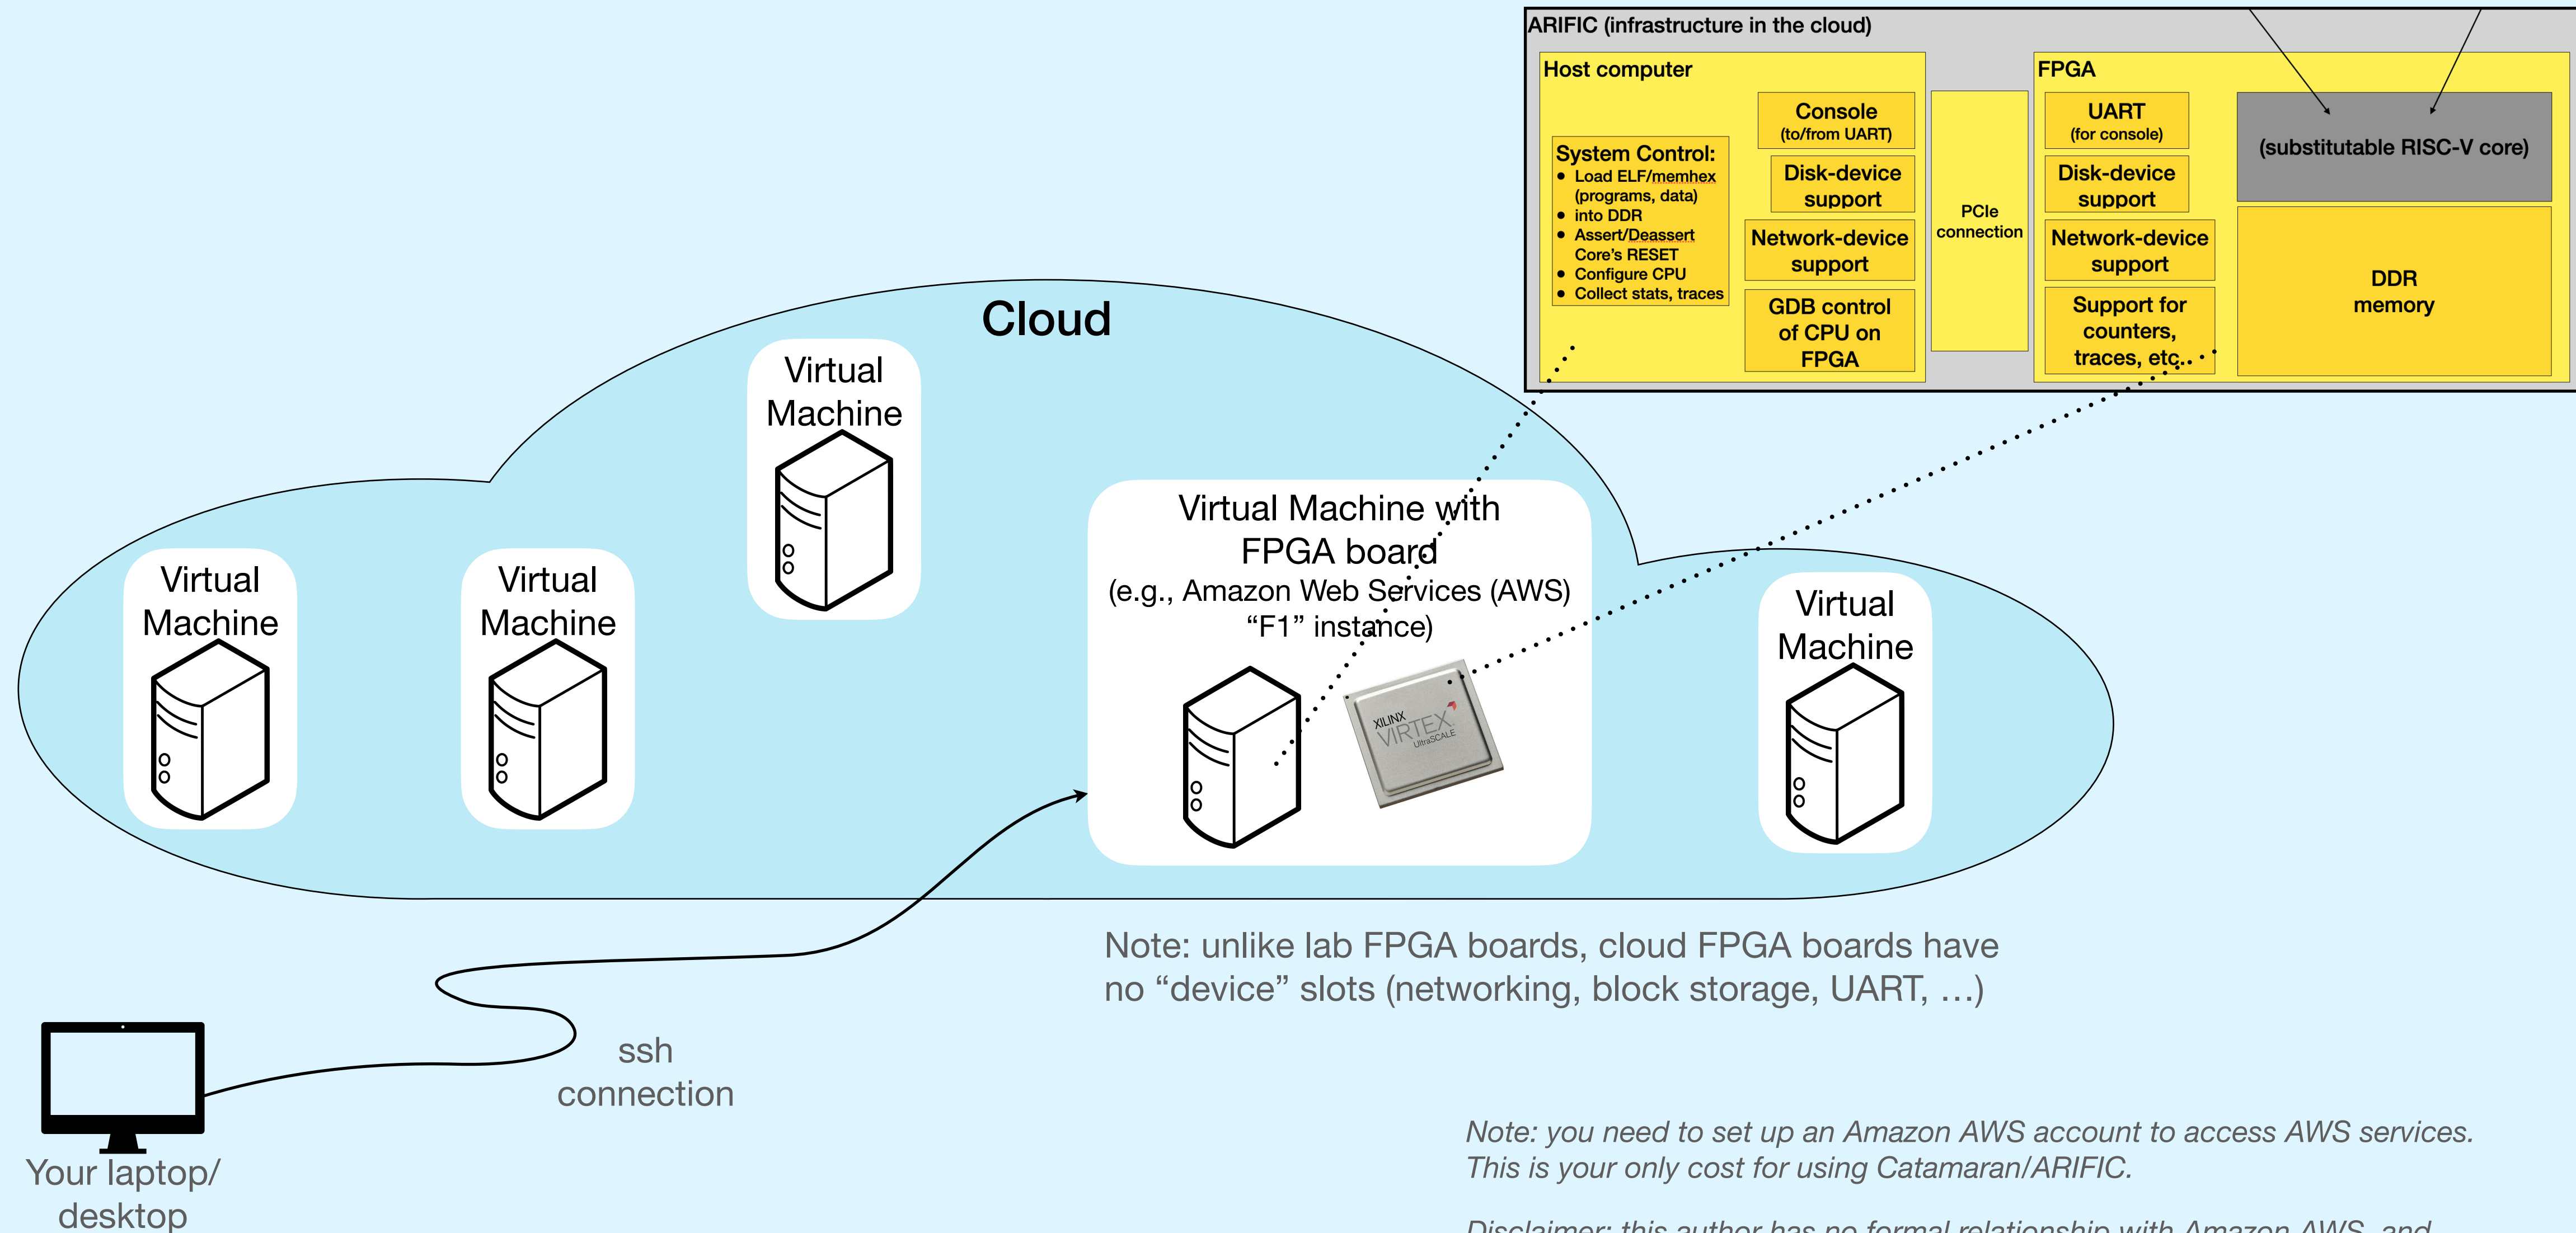
Network-device
support

Support for
counters,
traces, etc.

(substitutable RISC-V core)

DDR
memory

What are “FPGAs in the cloud”?



Note: unlike lab FPGA boards, cloud FPGA boards have no “device” slots (networking, block storage, UART, ...)

Note: you need to set up an Amazon AWS account to access AWS services. This is your only cost for using Catamaran/ARIFIC.

Disclaimer: this author has no formal relationship with Amazon AWS, and is merely an ordinary user of their services.

The rest of this tutorial

- Basics of Amazon AWS virtual machines, setting up an account, etc.
 - Skim during tutorial; details are for your future reference
- Demo/description of what you can do with Catamaran/ARIFIC
 - Run ISA tests
 - Cross-compile and run bare-metal (no OS) C programs
 - Run Linux, with networking and block devices
 - Cross-compile and run C programs under Linux
- Demo/description of plugging in your own RISC-V Core¹
 - Standard RTL-level interface for the core
 - Build a full-system simulation (using Verilator), and run it (can do on your local computer; does not need cloud)
 - Build a bitfile for AWS F1 FPGA, and run it
- Use GDB to control the RISC-V CPU on the FPGA

¹ “Your own RISC-V Core”


Options:

- Your own new CPU design
- Modify available CPU (many open-source)
 - microarchitecture change
 - new instruction
 - new CSRs (e.g., counters)
- Modify/replace memory system (caches, MMUs, PMPs, PTWs, ...)
- Add tightly-coupled or loosely coupled accelerator
- ... or other research idea ...

Basics of Amazon AWS

- *AMIs (virtual machines in the Amazon cloud)*
- Connecting to an AMI using “ssh”
- About AWS bitfiles

AMIs (virtual machines in Amazon cloud)



Amazon terminology for a virtual machine in their cloud:

“AMI” = “Amazon Machine Instance”

We also just say “instance” for an AMI/virtual machine

With an Amazon AWS account, you can “create” one or more AMIs for yourself (see Appendix for info on how to create an AMI).

For Catamaran/ARIFIC we usually create two AMIs:

- For *building* FPGA bitfiles: choose the free “FPGA Developer AMI” from AWS Marketplace, which is CentOS + Xilinx tools (Vivado) and licenses + aws-fpga SDK and HDK, running on “m6i.2xlarge” instance type.
- For *running* Catamaran/ARIFIC on FPGA: choose the free basic standard Ubuntu 22.04 AMI from AWS Marketplace, running on an “f1.2xlarge” instance type. “f1” instances have FPGAs attached.
- Costs:
 - Charges vary by instance type. Charges for “f1” instances (with FPGA) are somewhat higher, hence our choice above of a cheaper instance for our “build” AMI, which does not need an attached FPGA.
 - Charges are mostly for “running” instances; put your AMI into the “stopped” state when not in use.

Basics of Amazon AWS

- AMIs (virtual machines in the Amazon cloud)
- *DEMO: Connecting to an AMI using “ssh”*
- About AWS bitfiles

Connecting to your AMI with ssh (1/2)

In “Instances” dashboard ...

Select your instance

Use “Instance state” drop-down menu to start/stop selected instance

Make sure your AMI is “Running”

aws

Services

Search

[Option+S]

EC2

New EC2 Experience

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

Instances (1/3)

Find instance by attribute or tag (case-sensitive)

| | Name | Instance ID | Instance state | Instance type | Status check | Alarm status | Availability |
|-------------------------------------|----------------|---------------------|----------------|---------------|-------------------|--------------|--------------|
| <input checked="" type="checkbox"/> | RSN_AFI_Build2 | i-076651ea2ab3f3f46 | Running | m6i.2xlarge | 2/2 checks passed | No alarms | us-east-1a |
| <input type="checkbox"/> | RSN_AFI_Run | i-0f049486cb2592f76 | Stopped | f1.2xlarge | - | No alarms | us-east-1c |
| <input type="checkbox"/> | RSN_AWSteria | i-0b4c8df2758d2f0b9 | Stopped | f1.2xlarge | - | No alarms | us-east-1c |

Instance: i-076651ea2ab3f3f46 (RSN_AFI_Build2)

Details

Security

Networking

Storage

Status checks

Monitoring

Tags

Instance summary

Instance ID

i-076651ea2ab3f3f46 (RSN_AFI_Build2)

IPv6 address

-

Public IPv4 address

52.207.158.79 | open address

Instance state

Running

Private IPv4 addresses

172.31.47.209

Public IPv4 DNS

ec2-52-207-158-79.compute-1.amazonaws.com | open address

Note the IPv4 DNS address for the instance.
Hint: click to copy to clipboard

Connecting to your AMI with ssh (2/2)

In a terminal on your laptop/desktop, connect to your AMI:

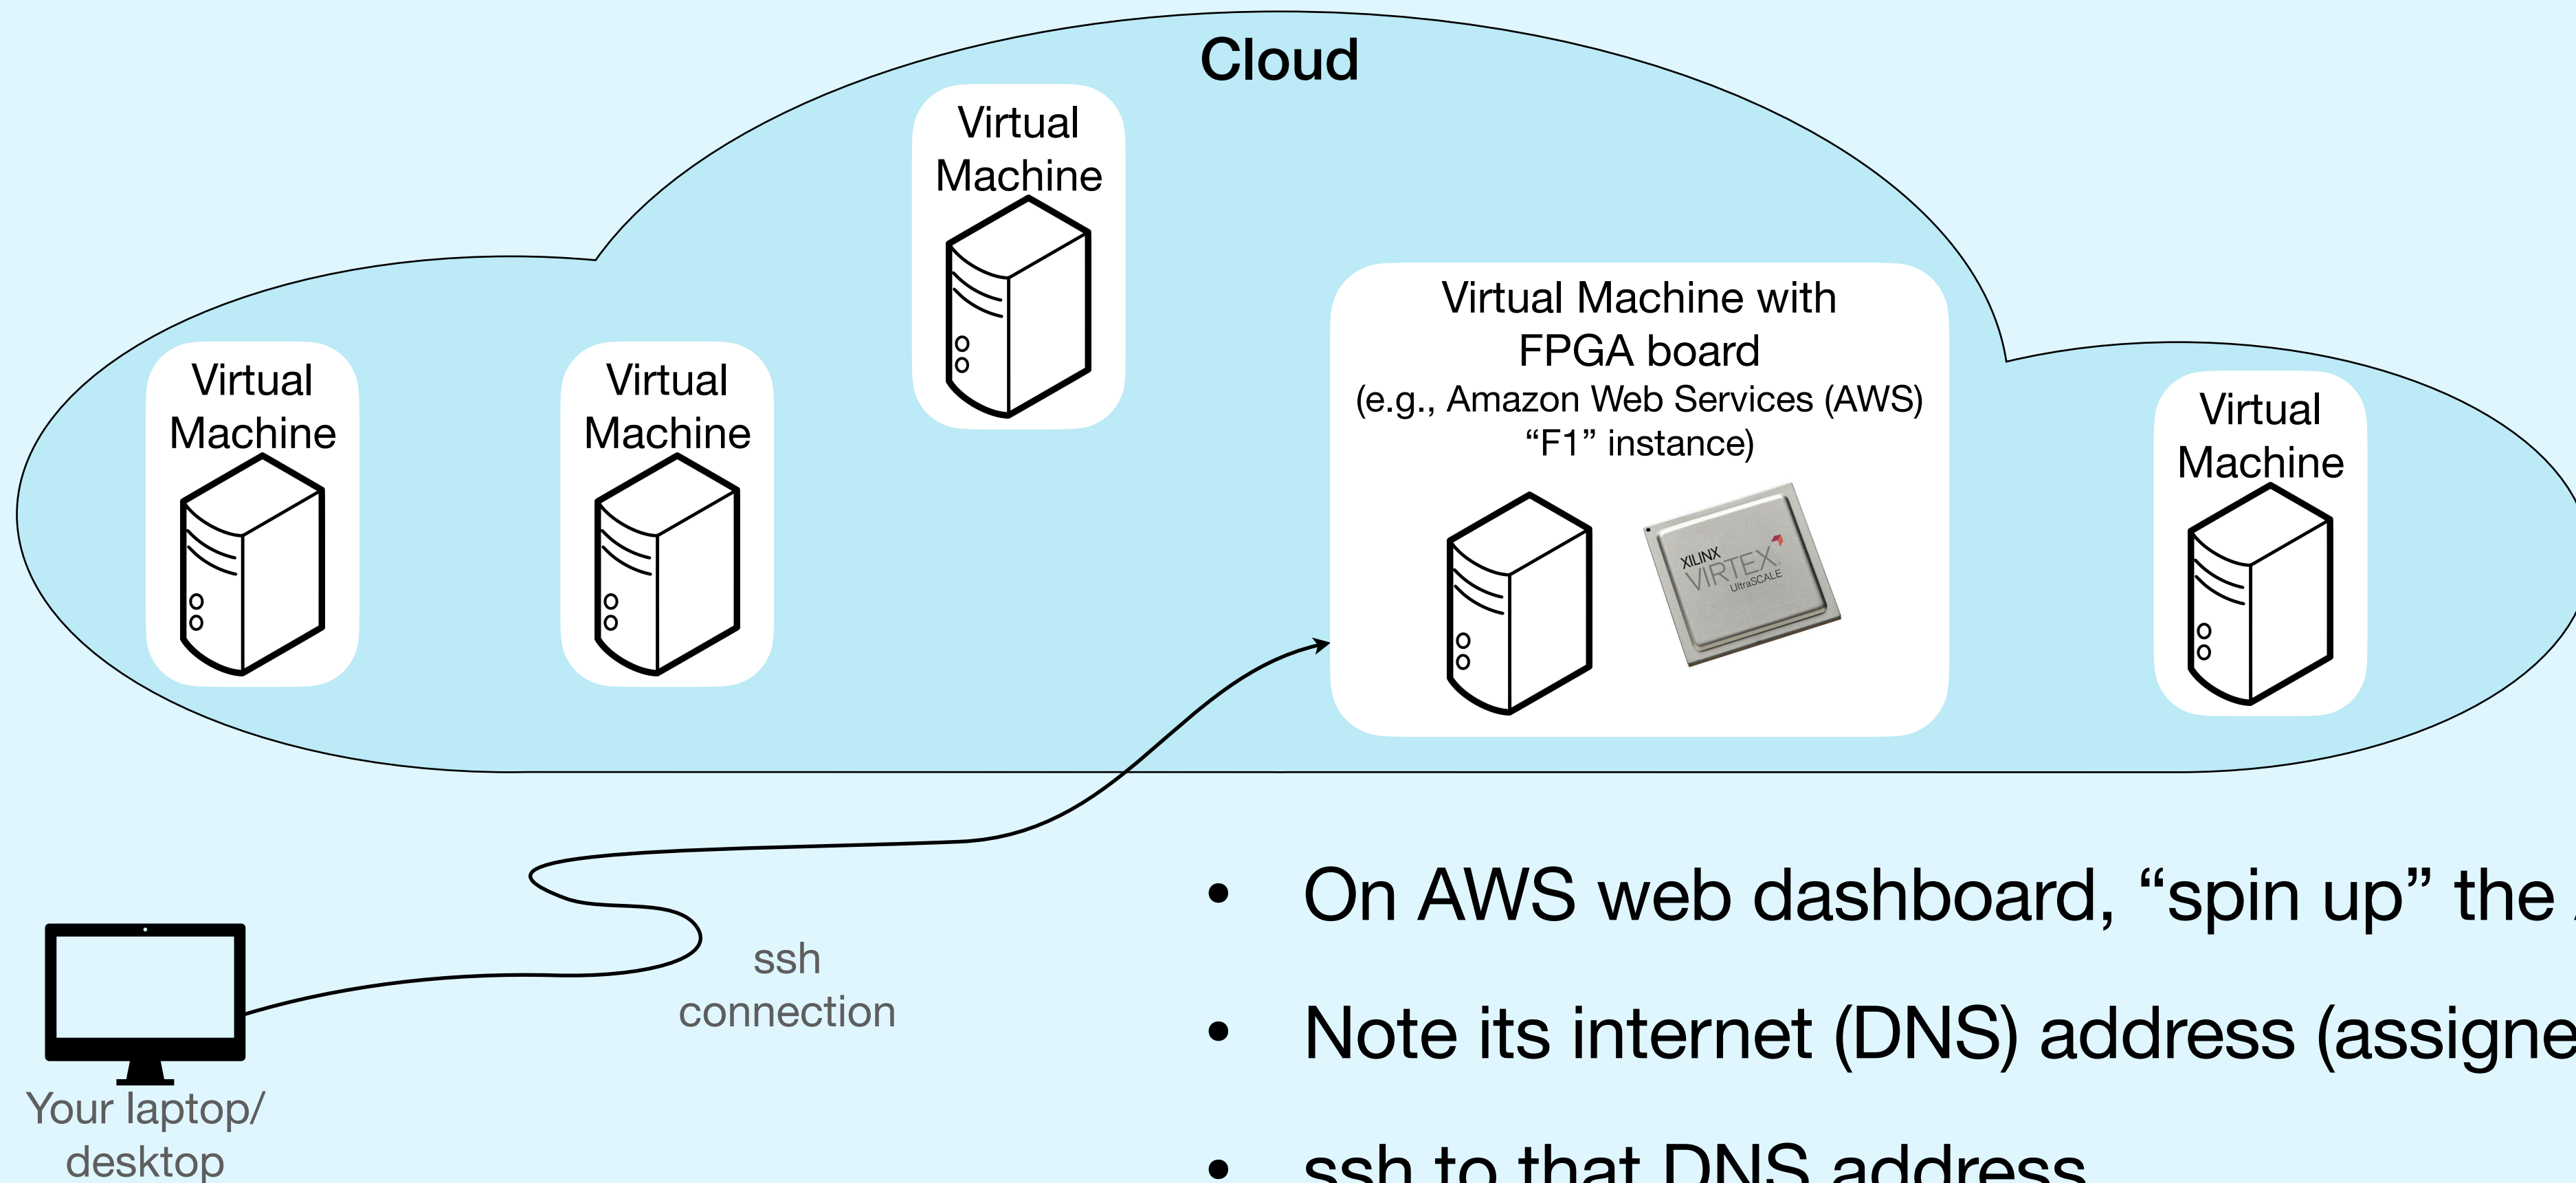
```
$ ssh -i ~/.ssh/MyPrivateKey.pem centos@ec2-52-207-158-79.compute-1.amazonaws.com
```

Note: you can copy files to and from your AMI using “scp” (which uses “ssh” under the covers)

```
$ scp -i ~/.ssh/MyPrivateKey.pem \  
    localfile \  
    centos@ec2-52-207-158-79.compute-1.amazonaws.com:~/remotefile
```

```
$ scp -i ~/.ssh/MyPrivateKey.pem \  
    centos@ec2-52-207-158-79.compute-1.amazonaws.com:~/remotefile \  
    localfile
```

DEMO: connect to your “run” AMI



- On AWS web dashboard, “spin up” the AMI into “running” state
- Note its internet (DNS) address (assigned by AWS during spin-up)
- ssh to that DNS address
- Tour of tutorial directory contents

Basics of Amazon AWS

- AMIs (virtual machines in the Amazon cloud)
- DEMO: Connecting to an AMI using “ssh”
- *About AWS bitfiles*

Basics of Amazon AWS

About AWS FPGA bitfiles (AFIs and AGFIs)

- Unlike “on-premises” FPGAs (on your lab bench/desktop), on AWS, FPGA bitfiles reside somewhere in the cloud (we have no direct access to the actual bitfile).
- Each bitfile has unique ids:
 - AFI (Amazon FPGA Image). These are unique within an AWS geographic region.
 - AGFI (...Global...). These are unique across all AWS geographic regions.
- See later section of this tutorial re. creating a bitfile and obtaining its AFI and AGFI
- AWS makes available software that, given an AGFI, will fetch the bitfile from somewhere in the cloud and load it into the FPGA in your instance.

The author does not know why AWS has two IDs for a bitfile, the AFI and AGFI, instead of just the AGFI?

Demo/description of what you can do with Catamaran/ARIFIC

- *DEMO: Cross-compile and run ISA tests*
- DEMO: Cross-compile and run bare-metal (no OS) C programs
- DEMO: Run Linux, with networking and block devices
- DEMO: Cross-compile and run C programs under Linux

Run ISA tests

What are “ISA tests”?

- RVI (RISC-V International, <https://riscv.org>) provides a standard set of tests for the RISC-V Instruction Set Architecture (ISA)
 - <https://github.com/riscv-software-src/riscv-tests>
- Consists of a number of assembly-language program files; each one focuses on testing one particular RISC-V instruction (opcode)
 - 236 tests for RV64 IMAFDC + MSU privileges + Sv39 Virtual Memory
 - 173 tests for RV32 IMAFDC + MSU privileges + Sv32 Virtual Memory
- Each program contains a number of sub-tests (each self-checking)
- Each program ends by writing to a particular “tohost” MMIO address
 - 0, if all sub-tests succeed (pass)
 - N, if sub-test N failed

Run ISA tests

Cross-compiling the RISC-V ISA tests

Please see Appendix for info on installing the (free) RISC-V GNU Toolchain (“gcc” and friends).

This tutorial repository contains, in the dir `ISA_Tests/`

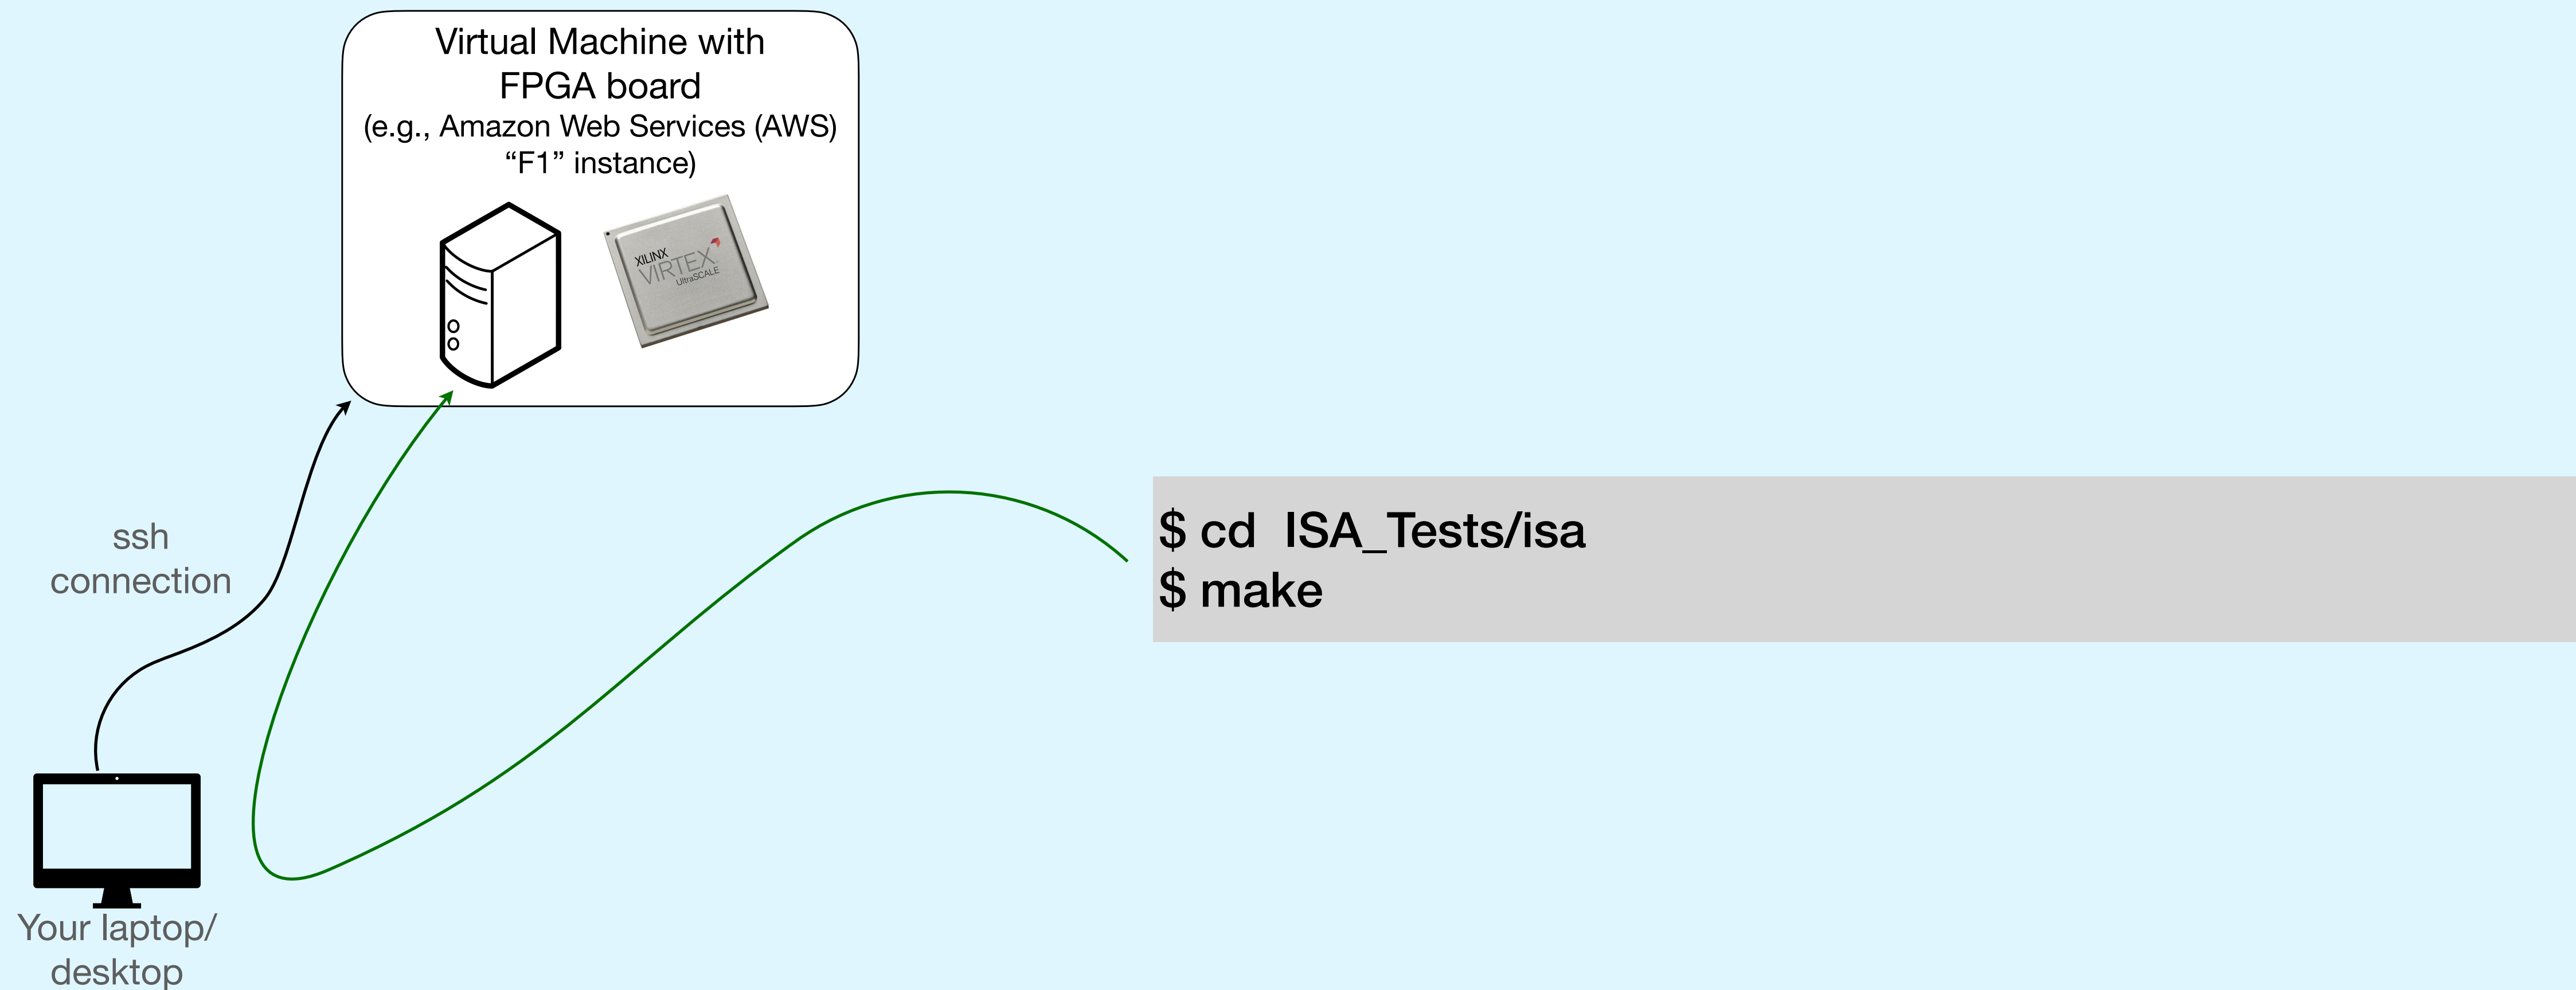
- A copy of the RISC-V ISA tests source code in `ISA_Tests/isa/` (see GitHub URL on previous slide)
- A Makefile and linker scripts to cross-compile the RISC-V ISA tests for RV64GC+MSU+Sv39
- Pre-built results of cross-compiling: set of RISC-V ELF files for the tests in `ISA_Tests/isa/elfs`

Note: our linker script sets

- Catamaran/ARIFIC program-start address = `0x_8000_0000`
- “tohost” MMIO address = `0x_6fff_0010`
- See `ISA_Tests/README.txt` for details

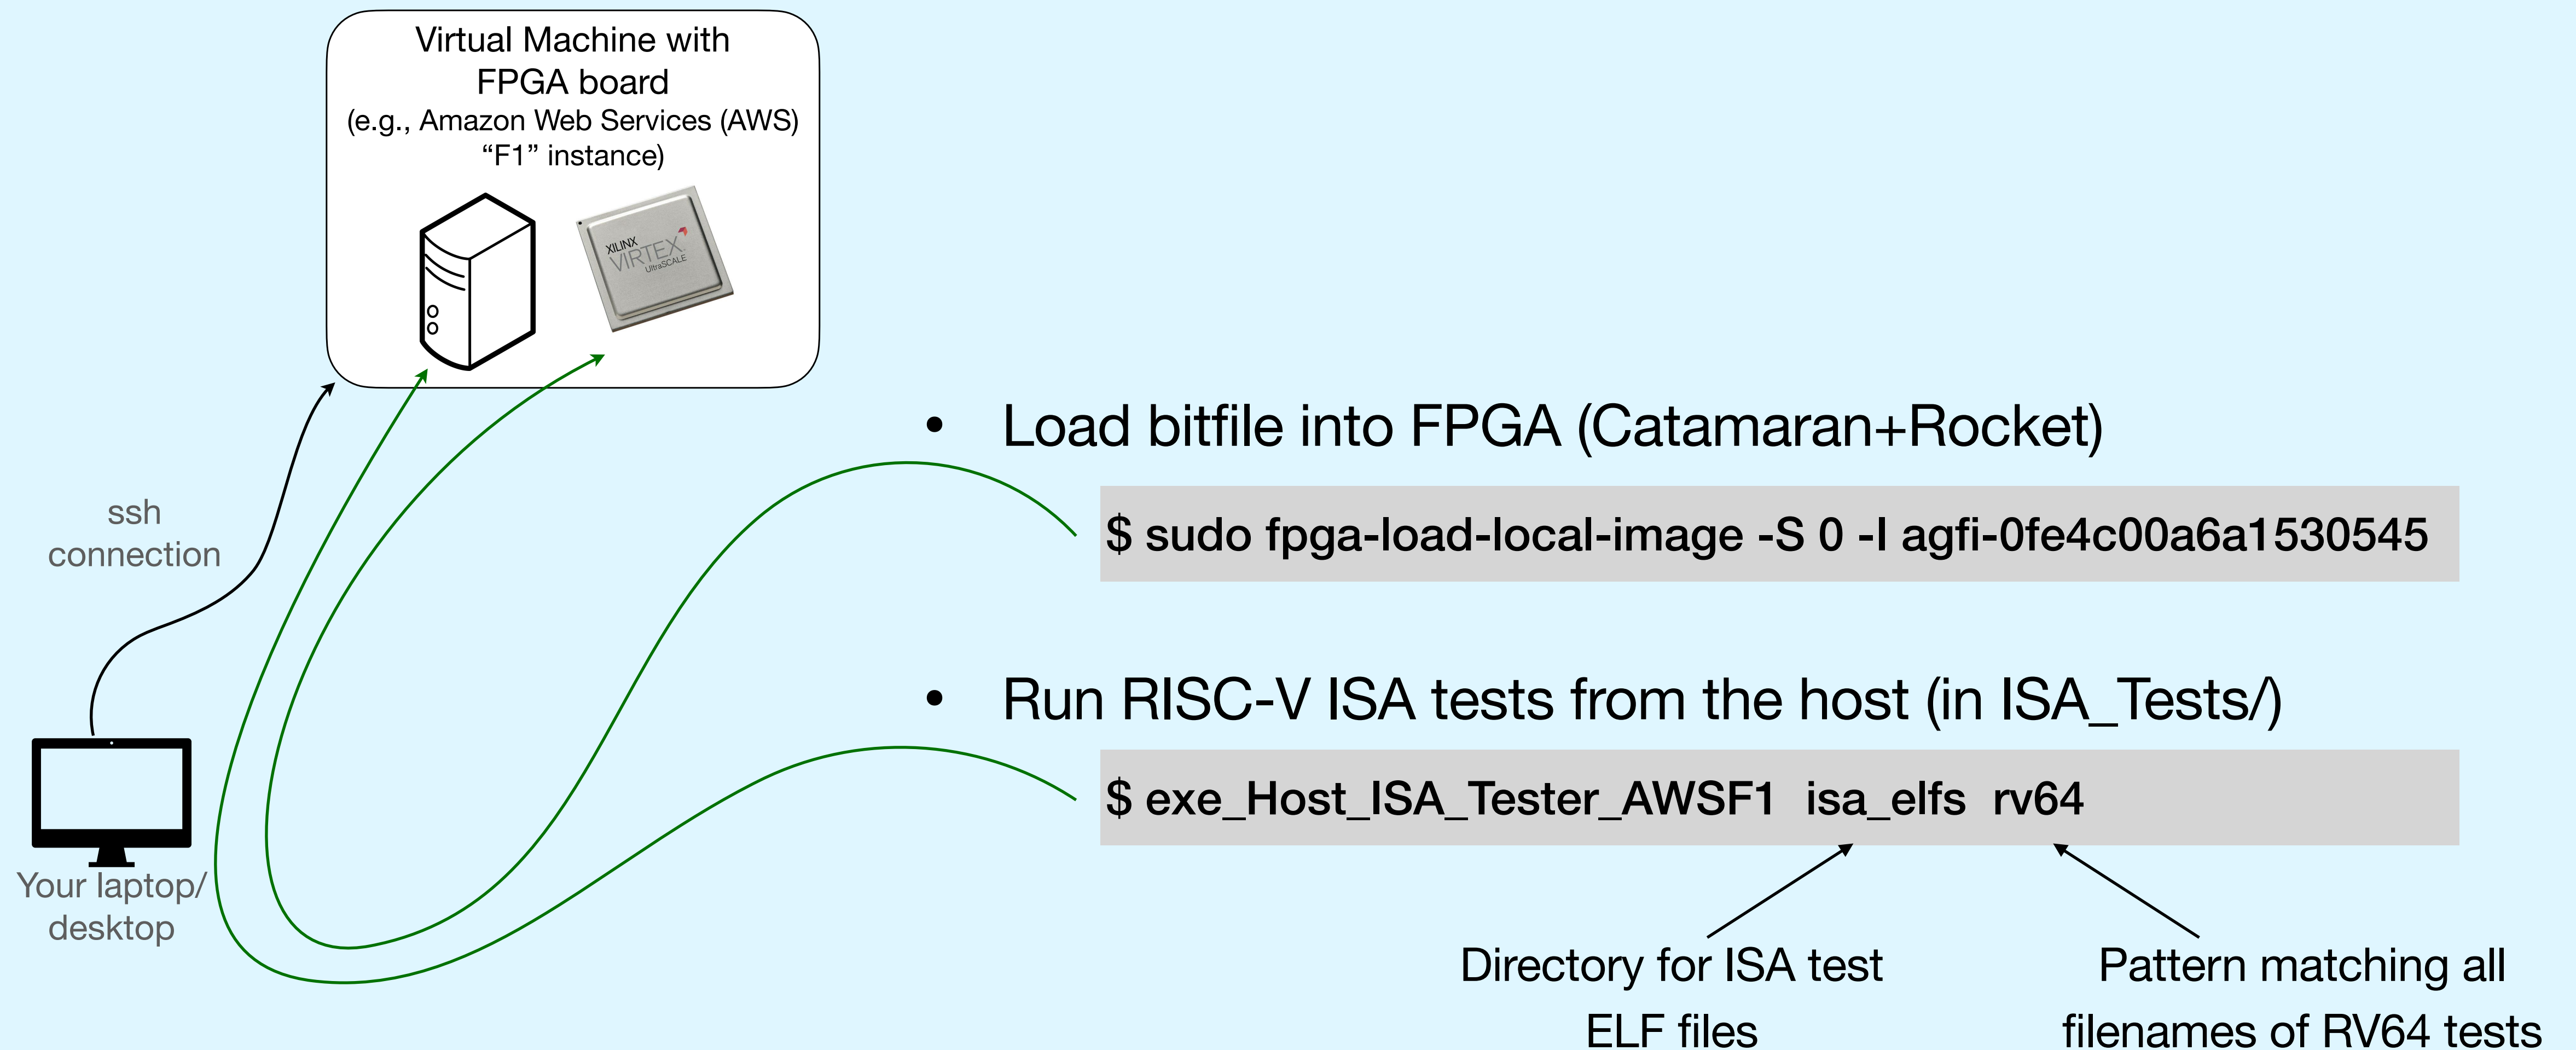
Run ISA tests

Demo: cross-compile RISC-V ISA tests



Run ISA tests

Demo: run RISC-V ISA tests



Demo/description of what you can do with Catamaran/ARIFIC

- DEMO: Cross-compile and run ISA tests
- *DEMO: Cross-compile and run bare-metal (no OS) C programs*
- DEMO: Run Linux, with networking and block devices
- DEMO: Cross-compile and run C programs under Linux

Cross-compile and run bare-metal (no OS) C programs

What are “bare-metal” C programs?

- No system calls. “printf/putc/putchar”, “scanf/getc/getchar” are linked with libraries that directly interact with the UART (device that sends/receives character to terminal)
- No “exit”; typically end in an infinite idle loop

This tutorial repository contains

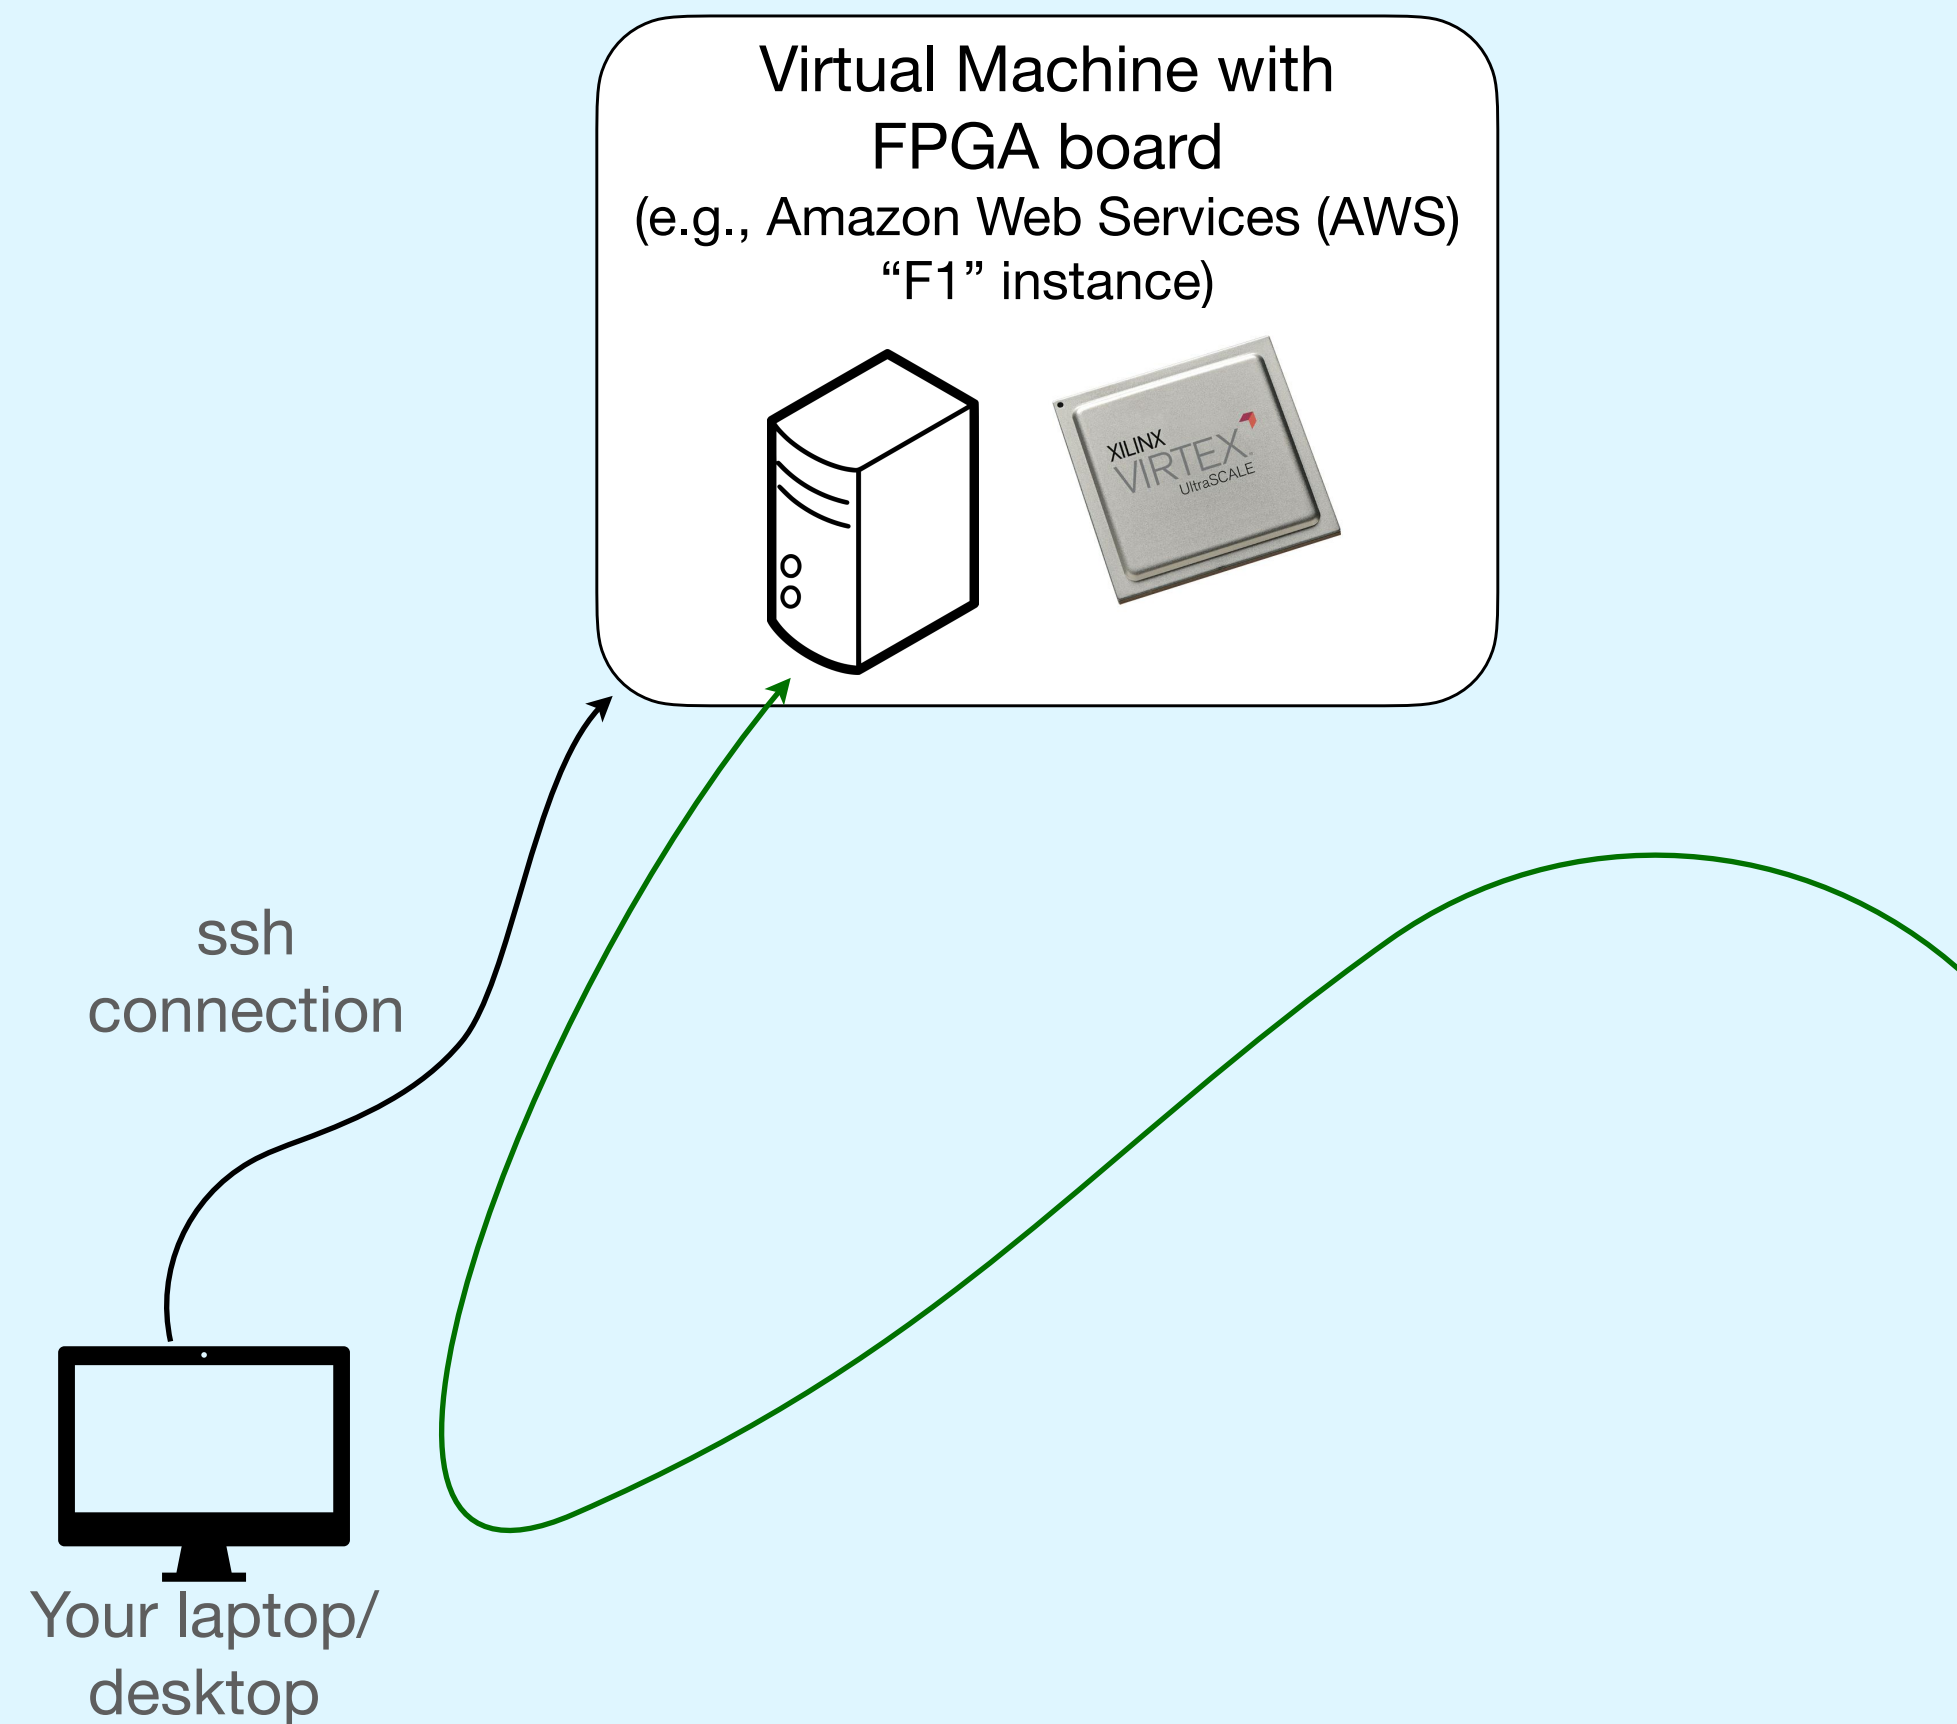
- A pre-built RISC-V GNU gcc cross-compiler (runs on Ubuntu, compiles for RV64GC a.k.a. RV64IMAFDC)
- Two small C programs: “Hello World!” and “cat” (echoes stdin to stdout)
- A Makefile and linker scripts to cross-compile the C programs for RV64GC
- Pre-built results of cross-compiling the examples: two RISC-V ELF files

Note: our linker script sets

- Catamaran/ARIFIC program-start address = 0x_8000_0000
- UART MMIO address = 0x_6010_0000

Run ISA tests

Demo: cross-compile bare-metal C programs



Classical “Hello World!” program

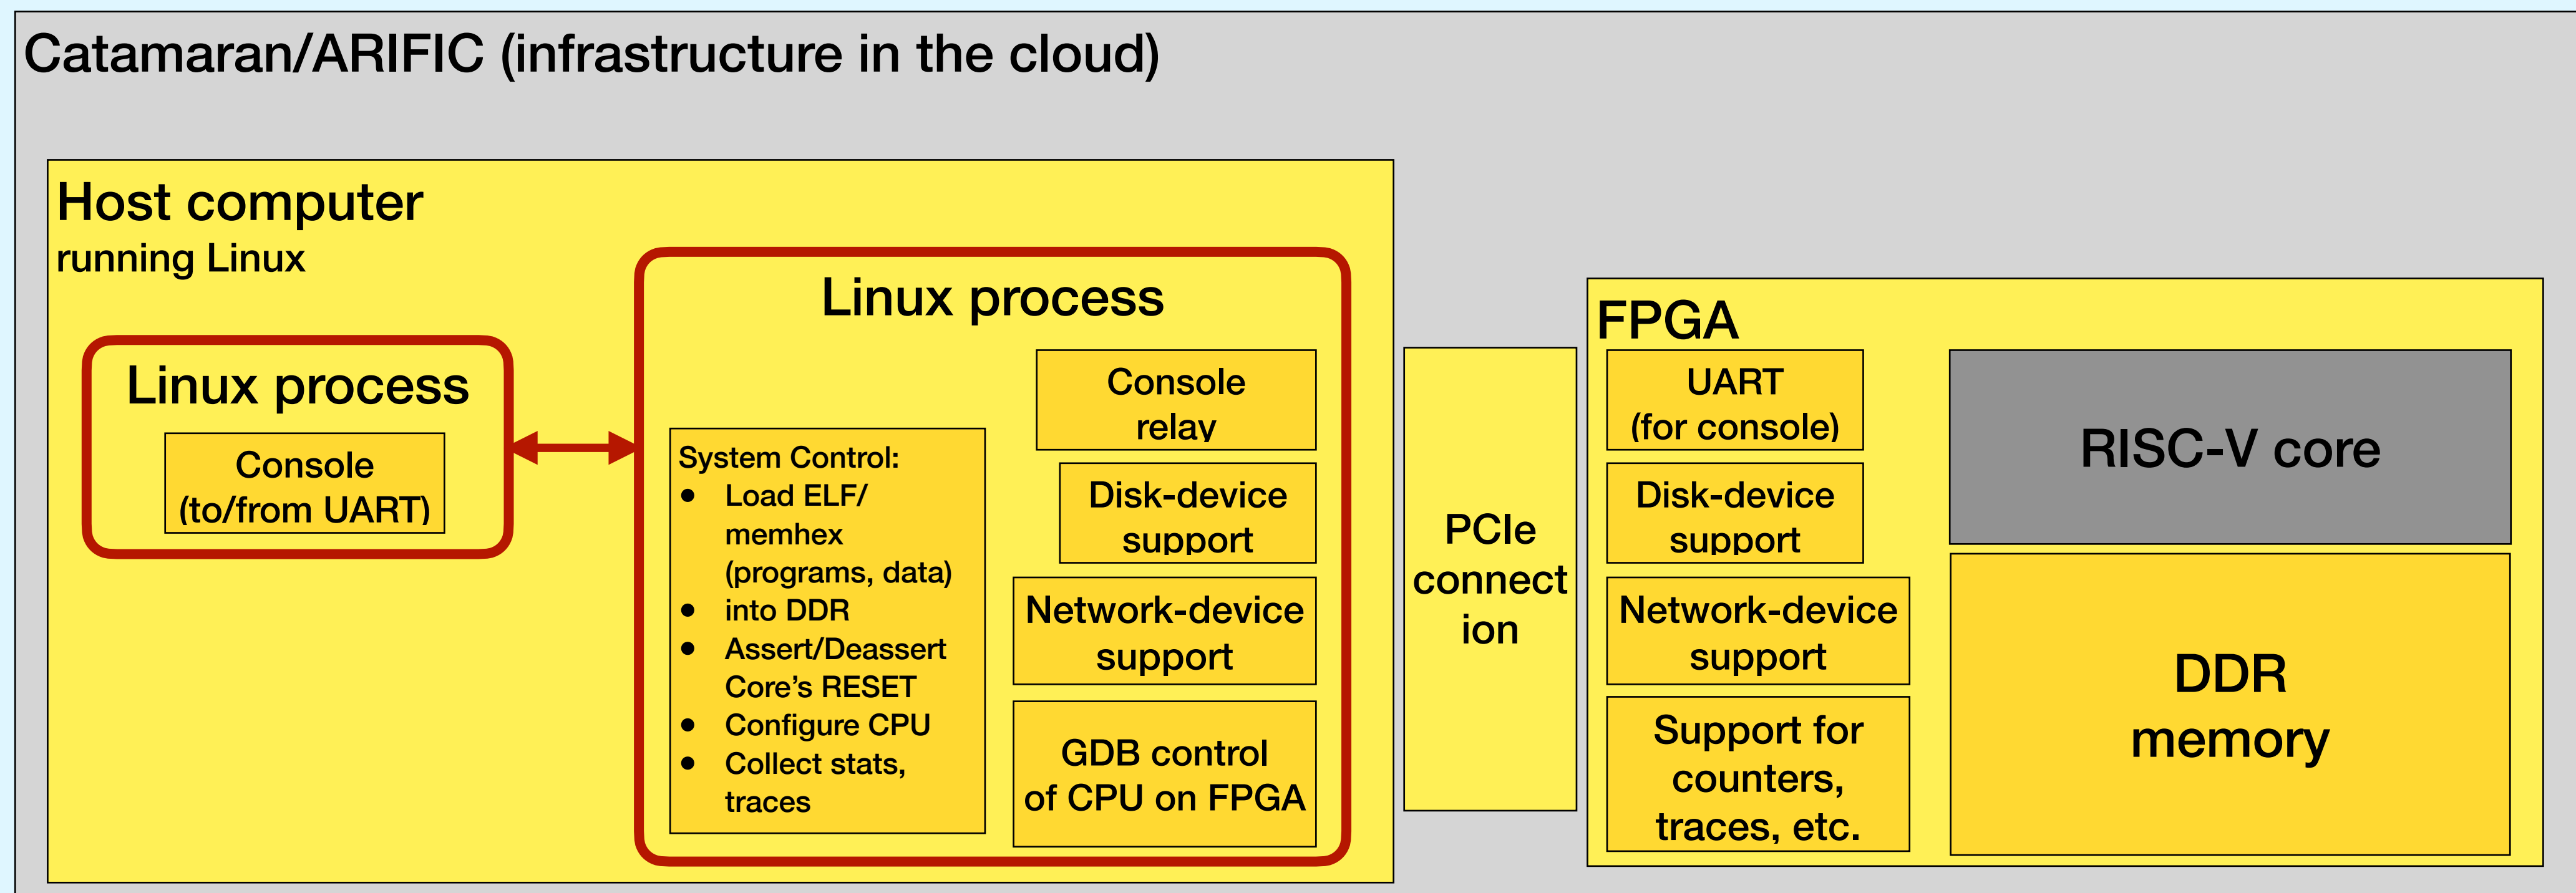
```
$ cd C_Examples/isa/hello  
$ make
```

Echo stdin to stdout

```
$ cd C_Examples/isa/cat  
$ make
```

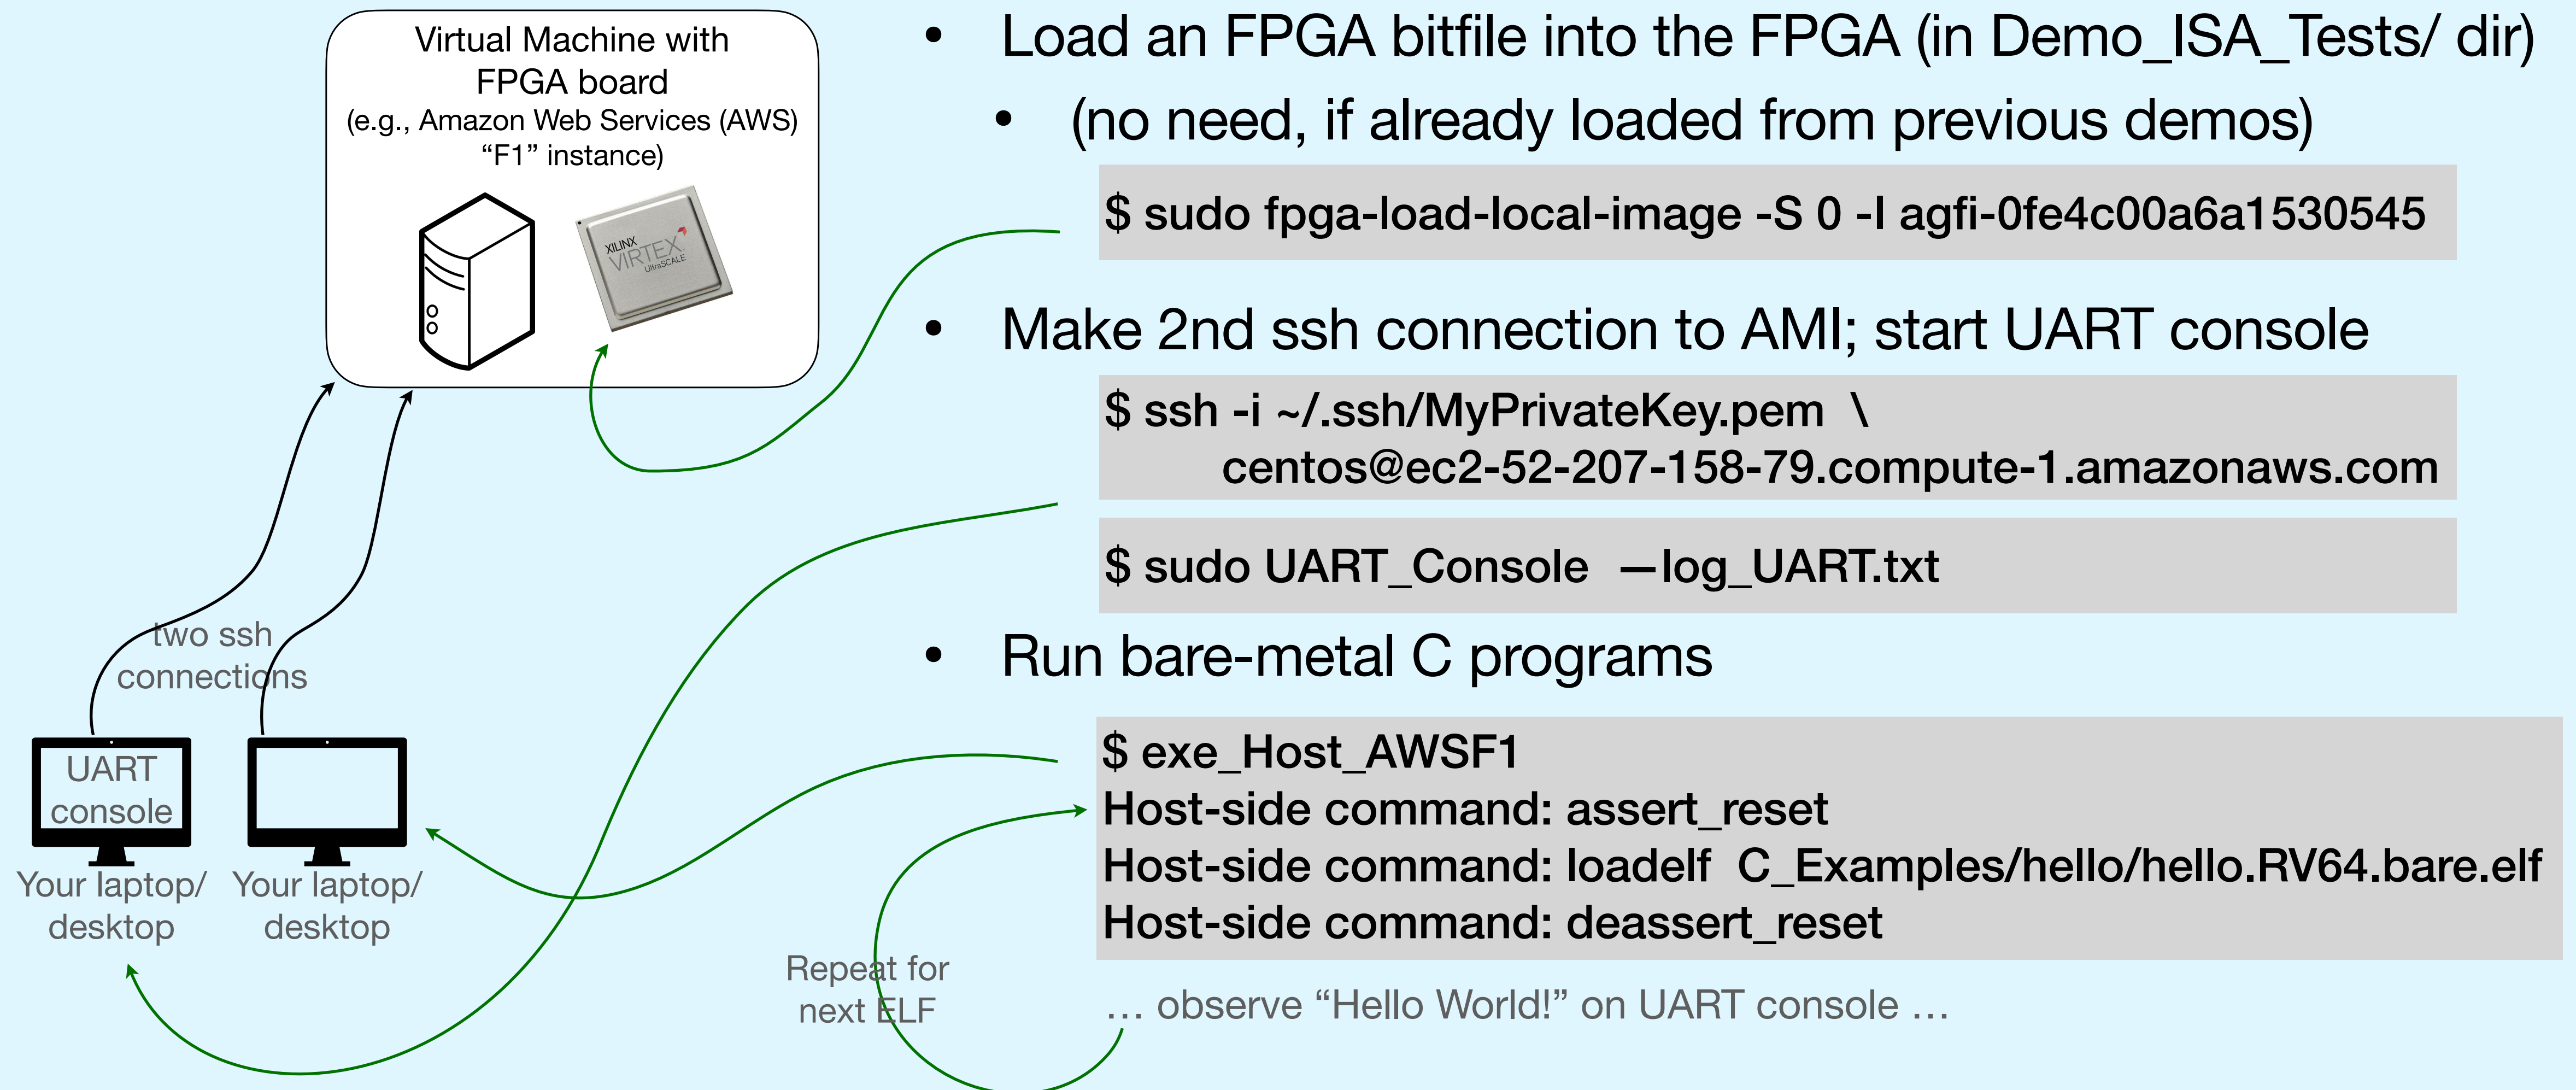
Cross-compile and run bare-metal (no OS) C programs

- Bare-metal C programs have no system calls. “printf/putc/putchar”, “scanf/getc/getchar” are linked with libraries that directly interact with the UART (hardware device that sends/receives characters)
- Catamaran/ARIFIC connects the UART to a terminal on the host computer



Cross-compile and run bare-metal (no OS) C programs

Demo: run bare-metal C programs



Demo/description of what you can do with Catamaran/ARIFIC

- DEMO: Cross-compile and run ISA tests
- DEMO: Cross-compile and run bare-metal (no OS) C programs
- *DEMO: Run Linux, with networking and block devices*
- DEMO: Cross-compile and run C programs under Linux

Demo/description of what you can do with Catamaran/ARIFIC

- DEMO: Cross-compile and run ISA tests
- DEMO: Cross-compile and run bare-metal (no OS) C programs
- DEMO: Run Linux, with networking and block devices
- *DEMO: Cross-compile and run C programs under Linux*


Demo/description of plugging in your own RISC-V Core

- The Catamaran/ARIFIC Core Interface (RTL)
- DEMO: Simulation flow (whole system, using Verilator)
- DEMO: FPGA flow on Amazon AWS

Plugging in your own RISC-V core: goal

Architecture Researcher's Focus

- CPU microarchitecture
- Memory systems (caches, MMUs, coherence, WMMs, ...)
- Accelerators



Catamaran/ARIFIC (infrastructure in the cloud)

Host computer running Linux

System Control:

- Load ELF/memhex (programs, data)
- into DDR
- Assert/Deassert Core's RESET
- Configure CPU
- Collect stats, traces

Console
(to/from UART)

Disk-device
support

Network-device
support

GDB control
of CPU on FPGA

PCIe
connection

FPGA

UART
(for console)

Disk-device
support

Network-device
support

Support for
counters,
traces, etc.

(substitutable RISC-V
core)

DDR
memory

- Assume you've created your own core¹
- Build a whole-system simulation executable of the FPGA-side, with your core substituted in place of the demo cores
- Build the FPGA bitfile of the FPGA-side, with your core substituted in place of the demo cores

¹ "Your own RISC-V Core"

Options:

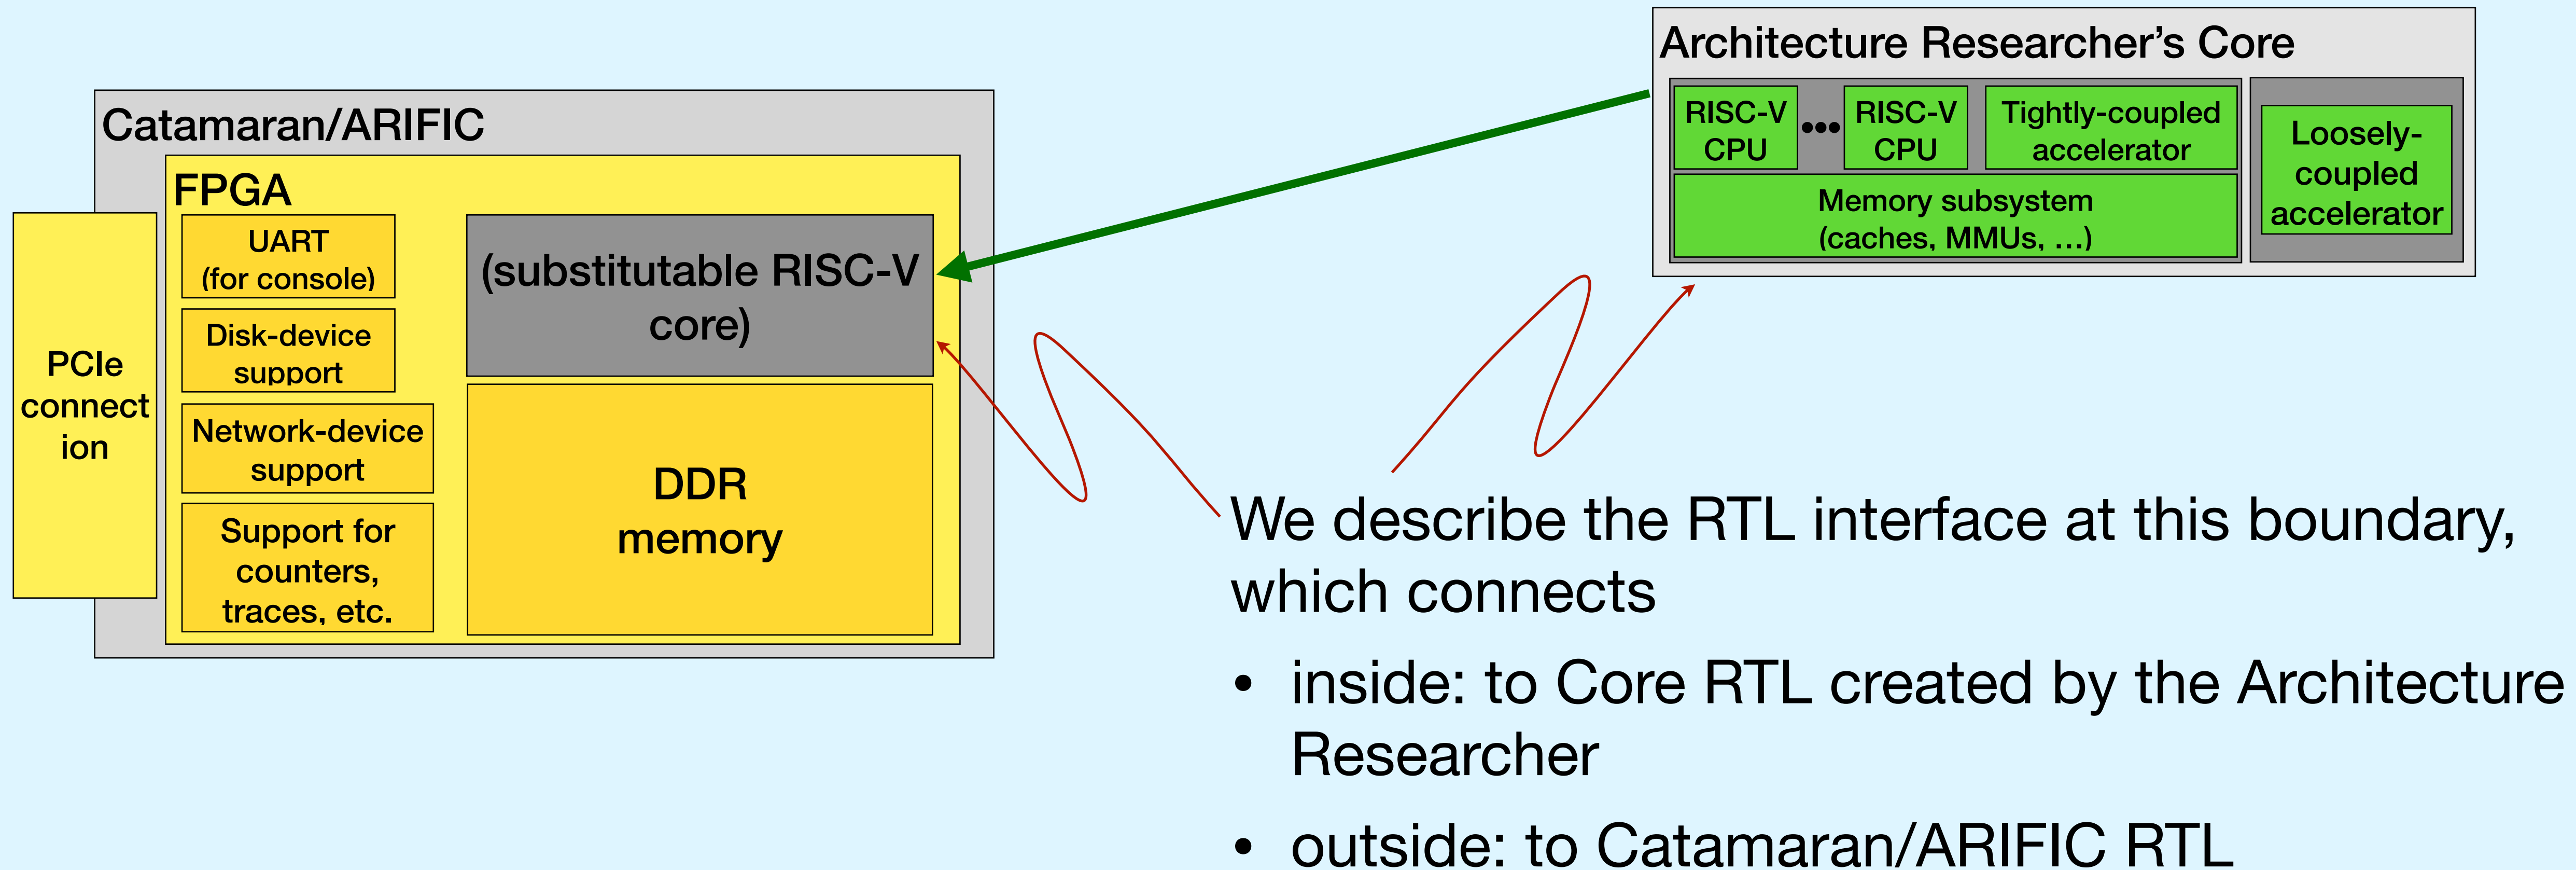
- Your own new CPU design
- Modify available CPU (many open-source)
 - microarchitecture change
 - new instruction
 - new CSRs (e.g., counters)
- Modify/replace memory system (caches, MMUs, PMPs, PTWs, ...)
- Add tightly-coupled or loosely coupled accelerator
- ... or other research idea ...

Demo/description of plugging in your own RISC-V Core

- *The Catamaran/ARIFIC Core Interface (RTL)*
- DEMO: Building a whole-system simulation executable
- DEMO: Building an FPGA bitfile for Amazon AWS

Plugging in your own RISC-V core

The Catamaran/ARIFIC Core Interface (RTL)

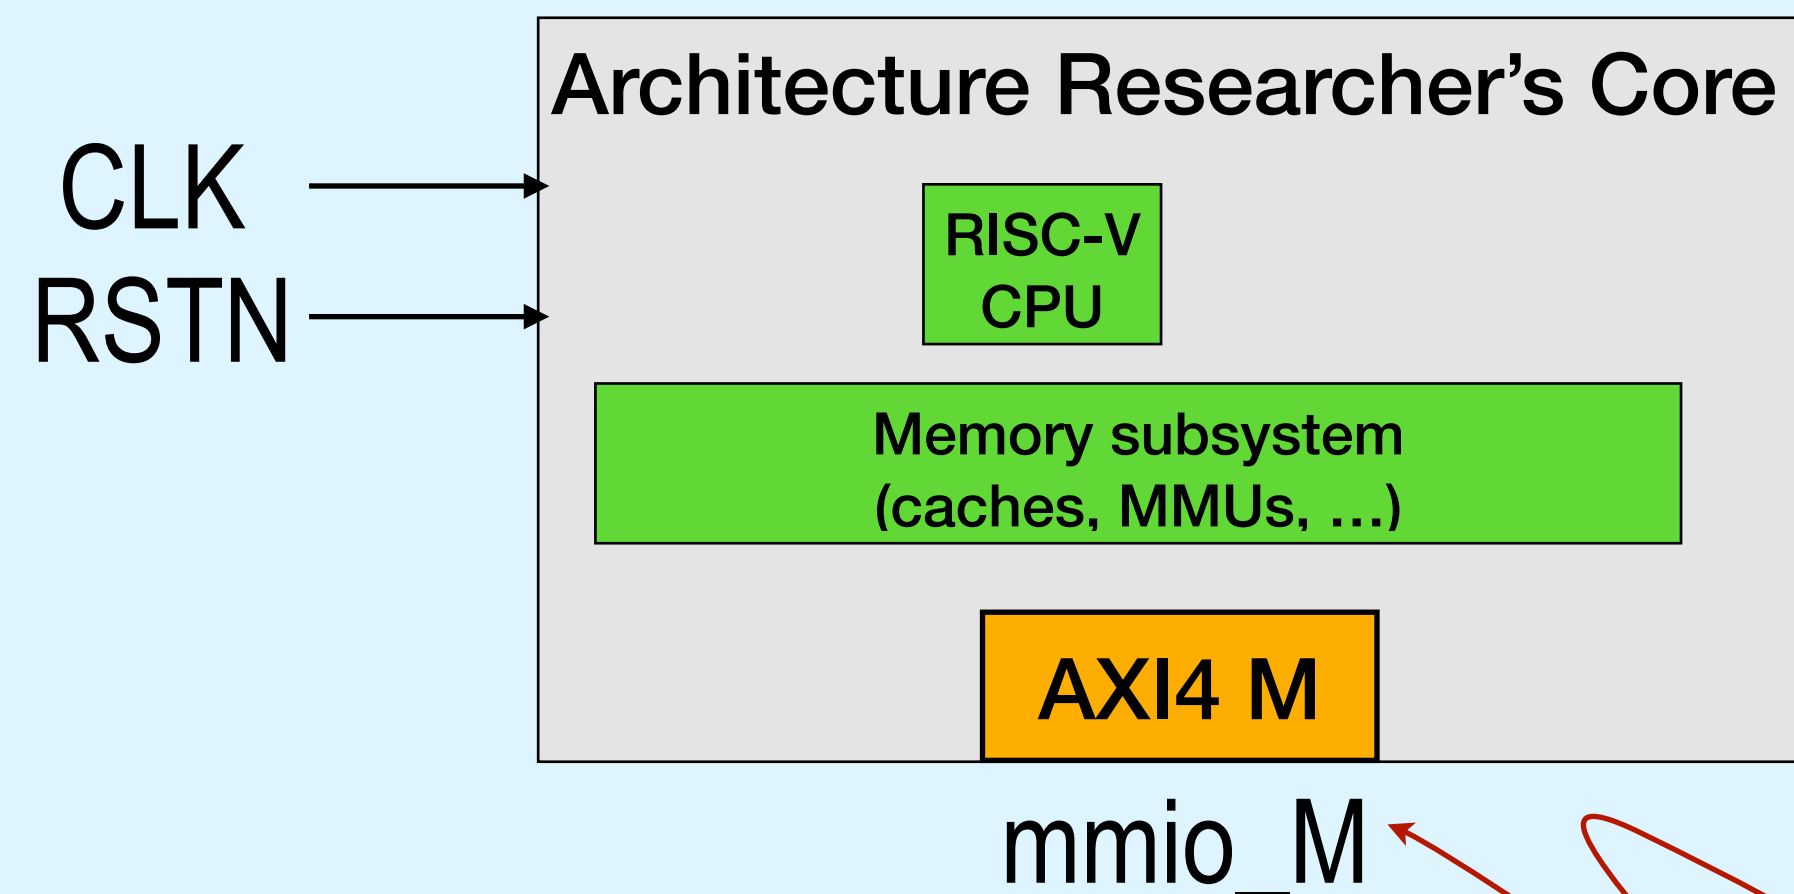


Plugging in your own RISC-V core

The Catamaran/ARIFIC Core Interface (RTL)

Demo: View the actual interface RTL in the tutorial repository

Minimal core



CPU can be RV32 or RV64,
from small (“embedded, IoT”)
to large (“application”, “server”).

AXI4 “manager” interface
64-bit address, 64-bit data,
connects to DDR memory,
UART, other devices


CLK can connect to one of six available clocks provided by Catamaran/ARIFIC, ranging from about 15 MHz to 250 MHz, depending on the speed/timing requirements of the core.

Plugging in your own RISC-V core

The Catamaran/ARIFIC Core Interface (RTL)

Demo: View the actual interface RTL in the tutorial repository

More interfaces




- In some cores, mem_M transfers entire cache lines at the back-end of an L2 cache
- In some cores, dma_S connects to a cache-coherent port in the memory subsystem

Plugging in your own RISC-V core

The Catamaran/ARIFIC Core Interface (RTL)

Demo: View the actual interface RTL in the tutorial repository




Demo/description of plugging in your own RISC-V Core

- The Catamaran/ARIFIC Core Interface (RTL)
- *DEMO: Building a whole-system simulation executable*
- DEMO: Building an FPGA bitfile for Amazon AWS

Plugging in your own RISC-V core

Building a whole-system simulation executable



Researcher's core can be written in any HDL, as long as you can generate RTL for it.

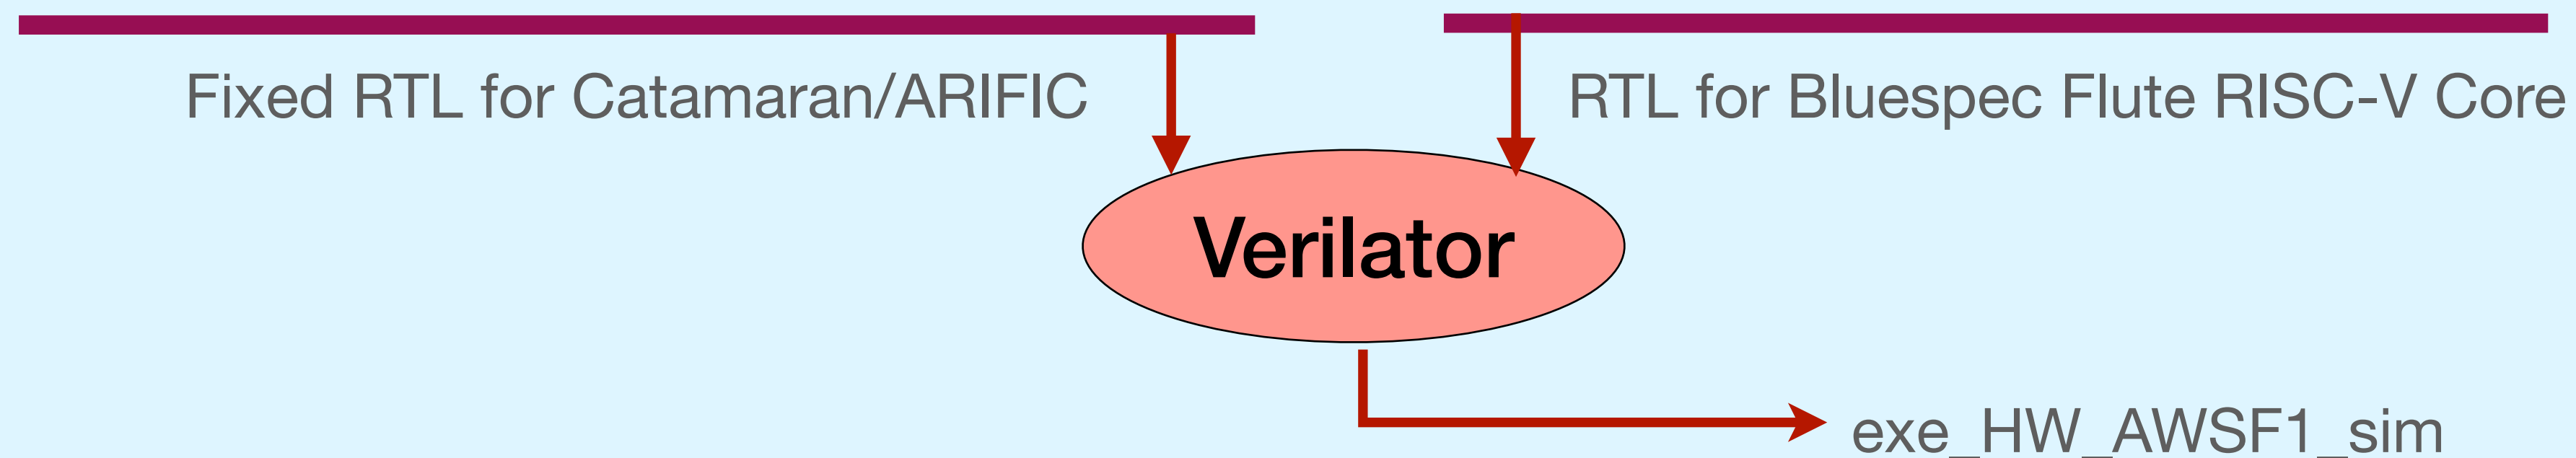
- E.g., Bluepec Flute is written in BSV
- E.g., Rocket is written in Chisel
- E.g., CVA6/Ariane is written in SystemVerilog

Plugging in your own RISC-V core

Demo: Building a whole-system simulation executable

This tutorial repository contains

- Fixed RTL for Catamaran/ARIFIC
- RTL for Bluespec Flute RISC-V Core (RV64GC_MSU_Sv39)
- Makefile for using Verilator to create a whole-system simulation executable
 - Note: Verilator is a free, open-source tool for creating RTL simulators
 - For installation, please see: <https://verilator.org/guide/latest/install.html>

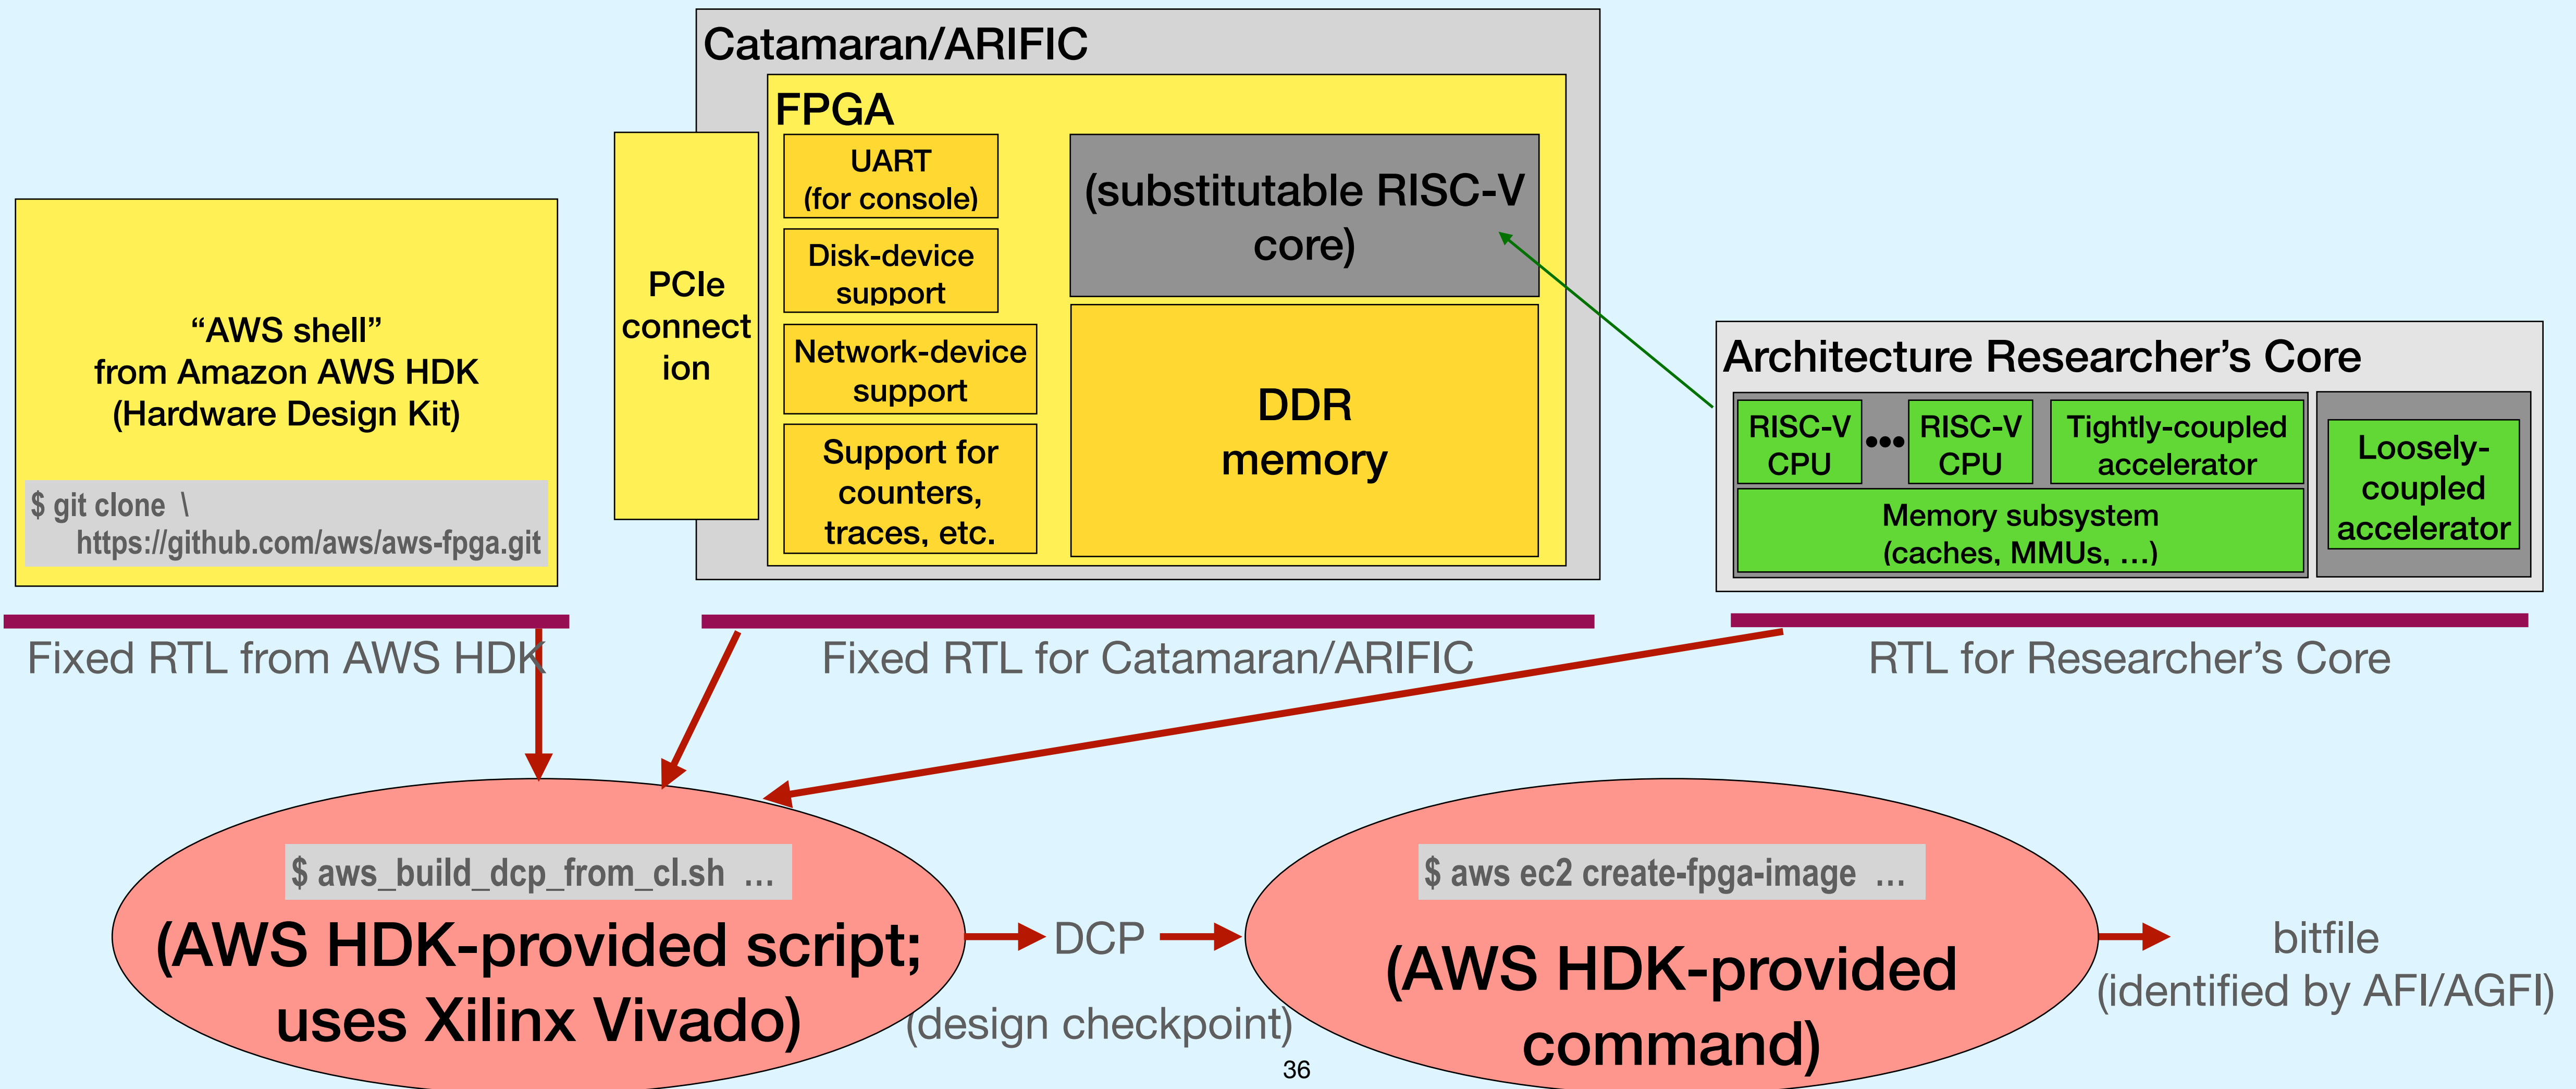


Demo/description of plugging in your own RISC-V Core

- The Catamaran/ARIFIC Core Interface (RTL)
- DEMO: Simulation flow (whole system, using Verilator)
- *DEMO: Building an FPGA bitfile for Amazon AWS*

Plugging in your own RISC-V core

Building an FPGA bitfile for Amazon AWS

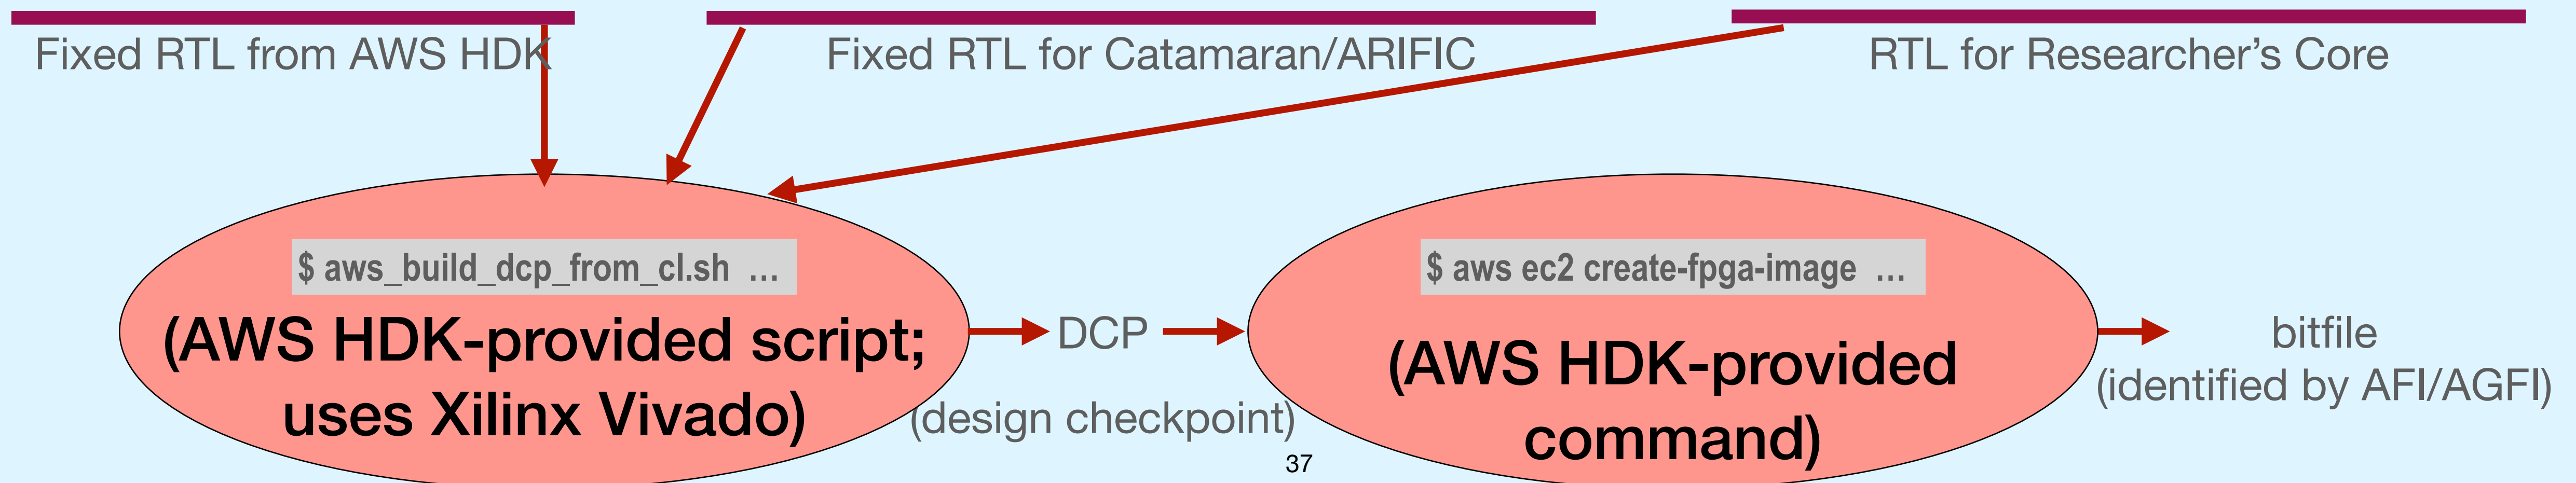


Plugging in your own RISC-V core

Demo: Building an FPGA bitfile for Amazon AWS

During the tutorial, we shall only demo the commands initiating these two steps, and not wait for the results


- The “build DCP” step can take several hours (e.g., 2h30m for building with CVA6 CPU)
- The “create FPGA image” step takes about 1 hour



**Use GDB to control the RISC-V
CPU on the FPGA**

Summary and Conclusion

Summary of tutorial



- Catamaran/ARIFIC provides the large and complex infrastructure that frees the Architecture Researcher to focus on their research into the CPU microarchitectures, memory systems, accelerators, etc.
- In this tutorial we showed the capabilities thereby immediately enabled for the Researcher:
 - Run on FPGA on Amazon AWS virtual machines
 - Run ISA tests
 - Run bare-metal C programs with console I/O
 - Run Linux and C programs-under-Linux with console I/O, networking and “disk” devices
 - GDB-debug bare-metal programs
 - Dump memory data (e.g., performance data collected during program execution)

End of Tutorial Slides; Reference Material follows

Thank you all for attending!

GitHub repo for this tutorial: https://github.com/rsnikhil/Tutorial_at_HPCA-29

Appendix: Reference Material

- *Installing the RISC-V Gnu Toolchain (gcc, as, ld, gdb, ...)*
- Opening an account on Amazon AWS
- Creating an AMI (or two) for Catamaran/ARIFIC

Installing the RISC-V Gnu Toolchain

- By “RISC-V Gnu Toolchain” we mean the familiar *gcc*, *as*, *ld*, *gdb*, *objdump*, ... etc.
- These are needed for compiling C, C++ and RISC-V Assembly Language programs into ELF binaries to be loaded and run on a RISC-V CPU.
- We install “cross-compilers” that run under Ubuntu, but generate RISC-V machine code.

Full details may be found at this URL; please follow the directions there.

<https://github.com/riscv/riscv-gnu-toolchain>

Hint: in the “configure” step, specify “medany” and “multilib”:

```
$ ./configure --prefix=<path_to_installation_dir> --with-cmodel=medany --enable-multilib
```

- “newlib” is for the compiler for bare-metal C RV32 and RV64 programs (no system calls)
- “multilib” produces a compiler that can produce both RV32 and RV64 code, even though it is named *riscv64-unknown-elf-gcc*


Remember to export the environment variable “RISCV” to point at the installation directory, and to place the *bin* directory in your `PATH`.

Appendix: Reference Material

- Installing the RISC-V Gnu Toolchain (gcc, as, ld, gdb, ...)
- *Opening an account on Amazon AWS*
- Creating an AMI (or two) for Catamaran/ARIFIC

Create your Amazon AWS account

- At: <https://aws.amazon.com>



The screenshot shows the AWS sign-in page. It includes the AWS logo, a 'Sign in' header, and two user type options: 'Root user' (selected) and 'IAM user'. Below these is a text input field for 'Root user email address' containing 'username@example.com', followed by a blue 'Next' button. A paragraph of terms and conditions follows. At the bottom, there is a 'New to AWS?' section with a 'Create a new AWS account' button. Two arrows originate from the right side: one labeled 'First time' pointing to the 'Next' button, and another labeled 'Subsequently' pointing to the 'Create a new AWS account' button.

First time

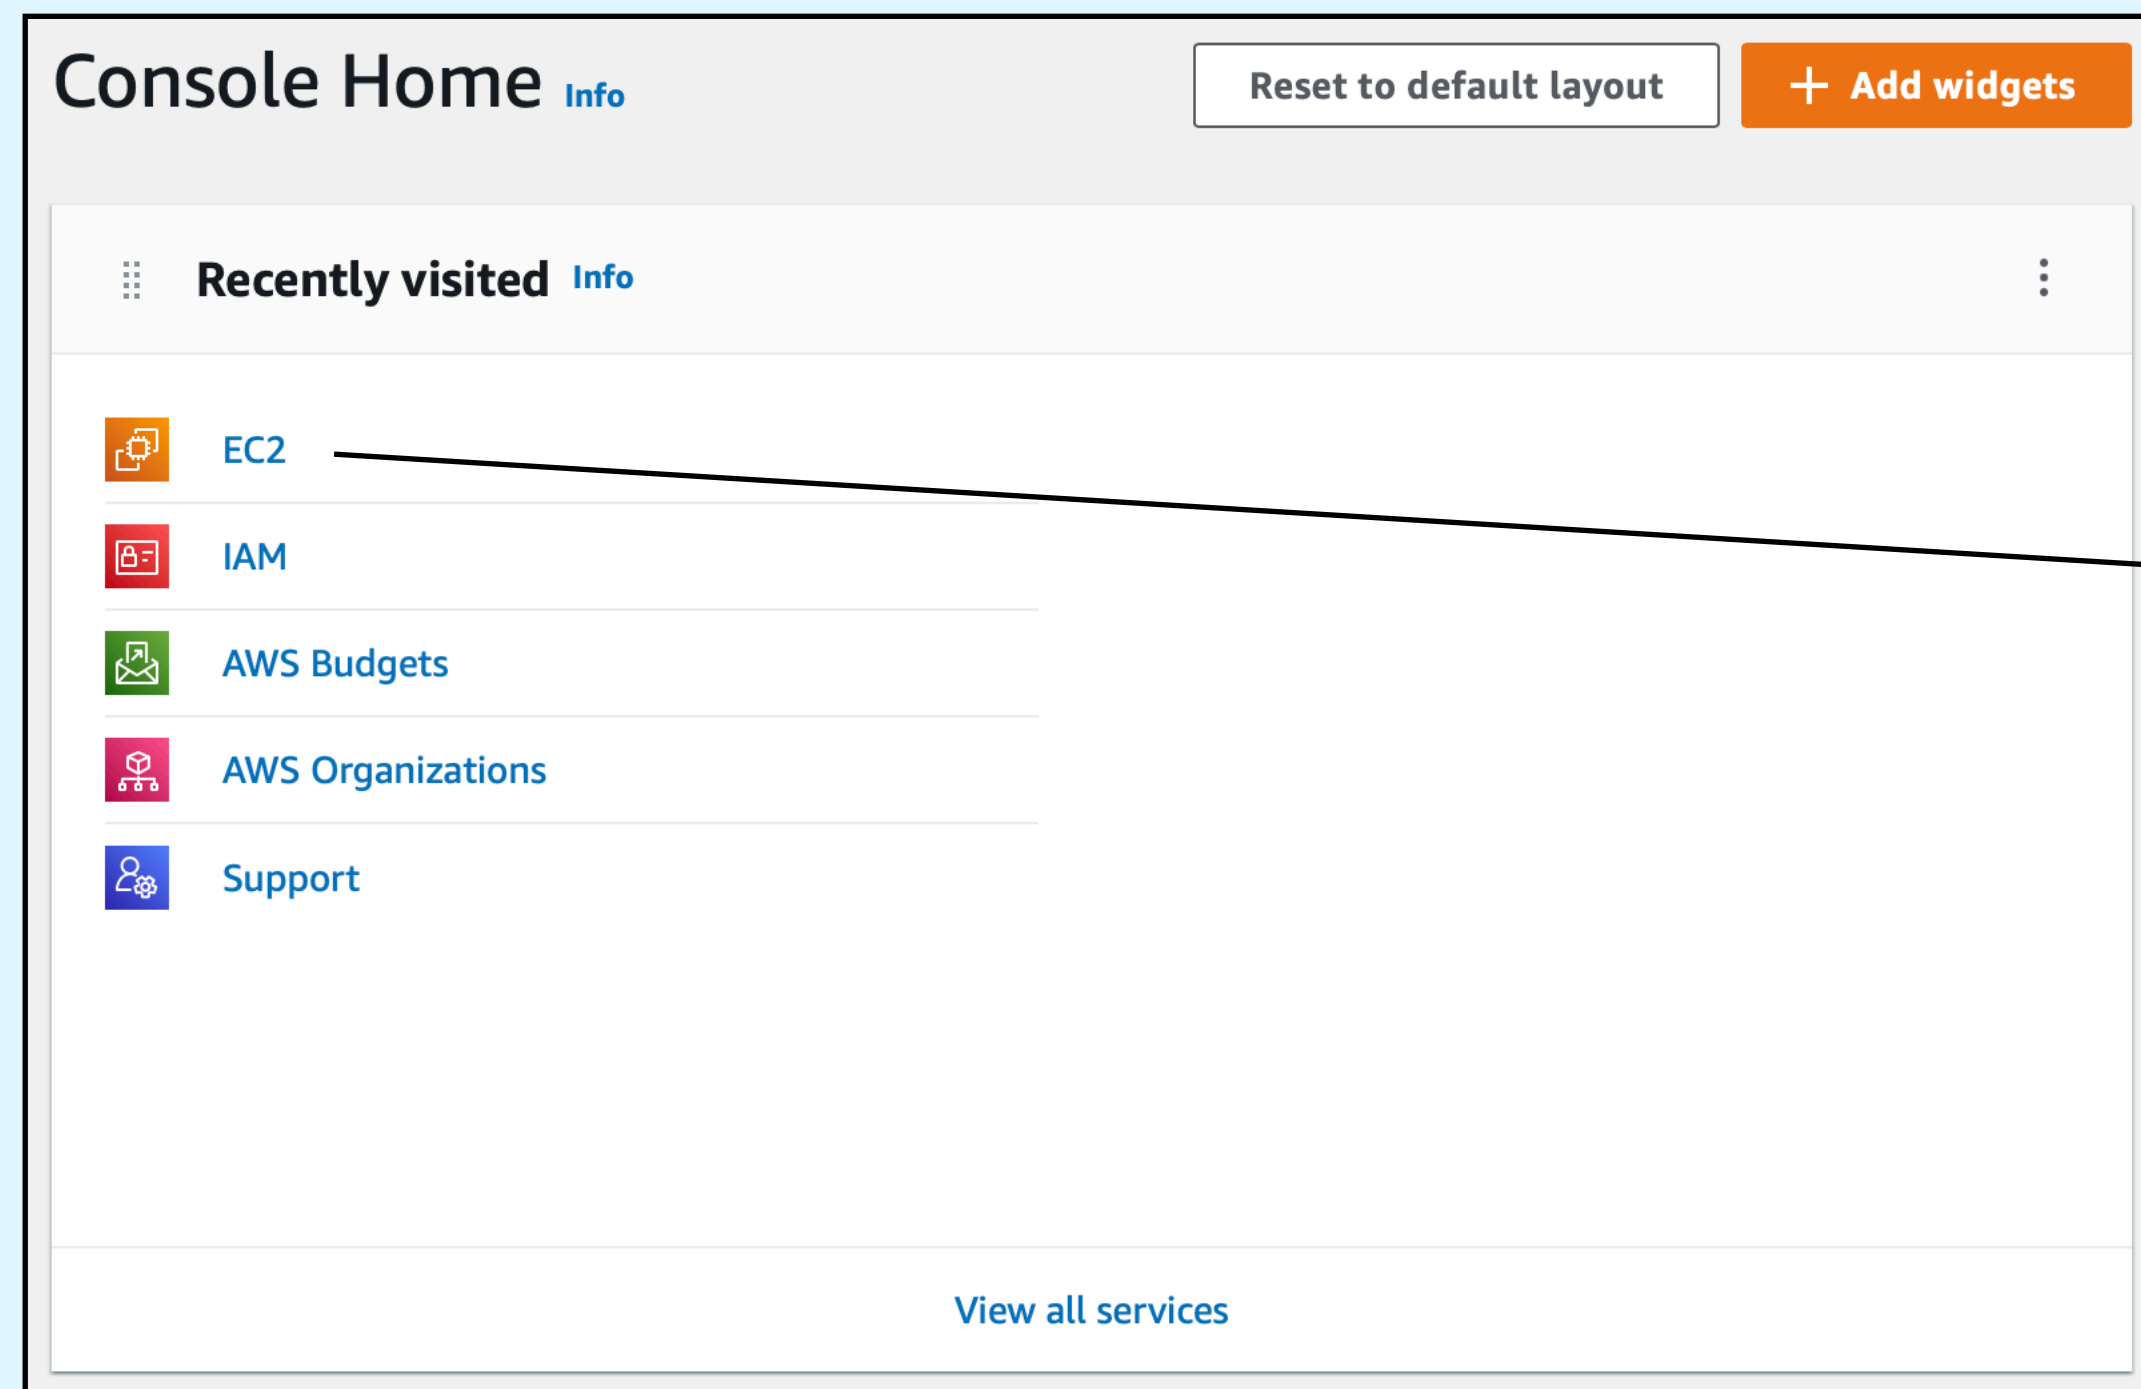
Subsequently

Amazon runs cloud farms in many parts of the world; you may be asked to select a region “near” you.
E.g., I use: ‘us-east-1 (N.Virginia)’
Caution: select a region that has support instances with FPGAs

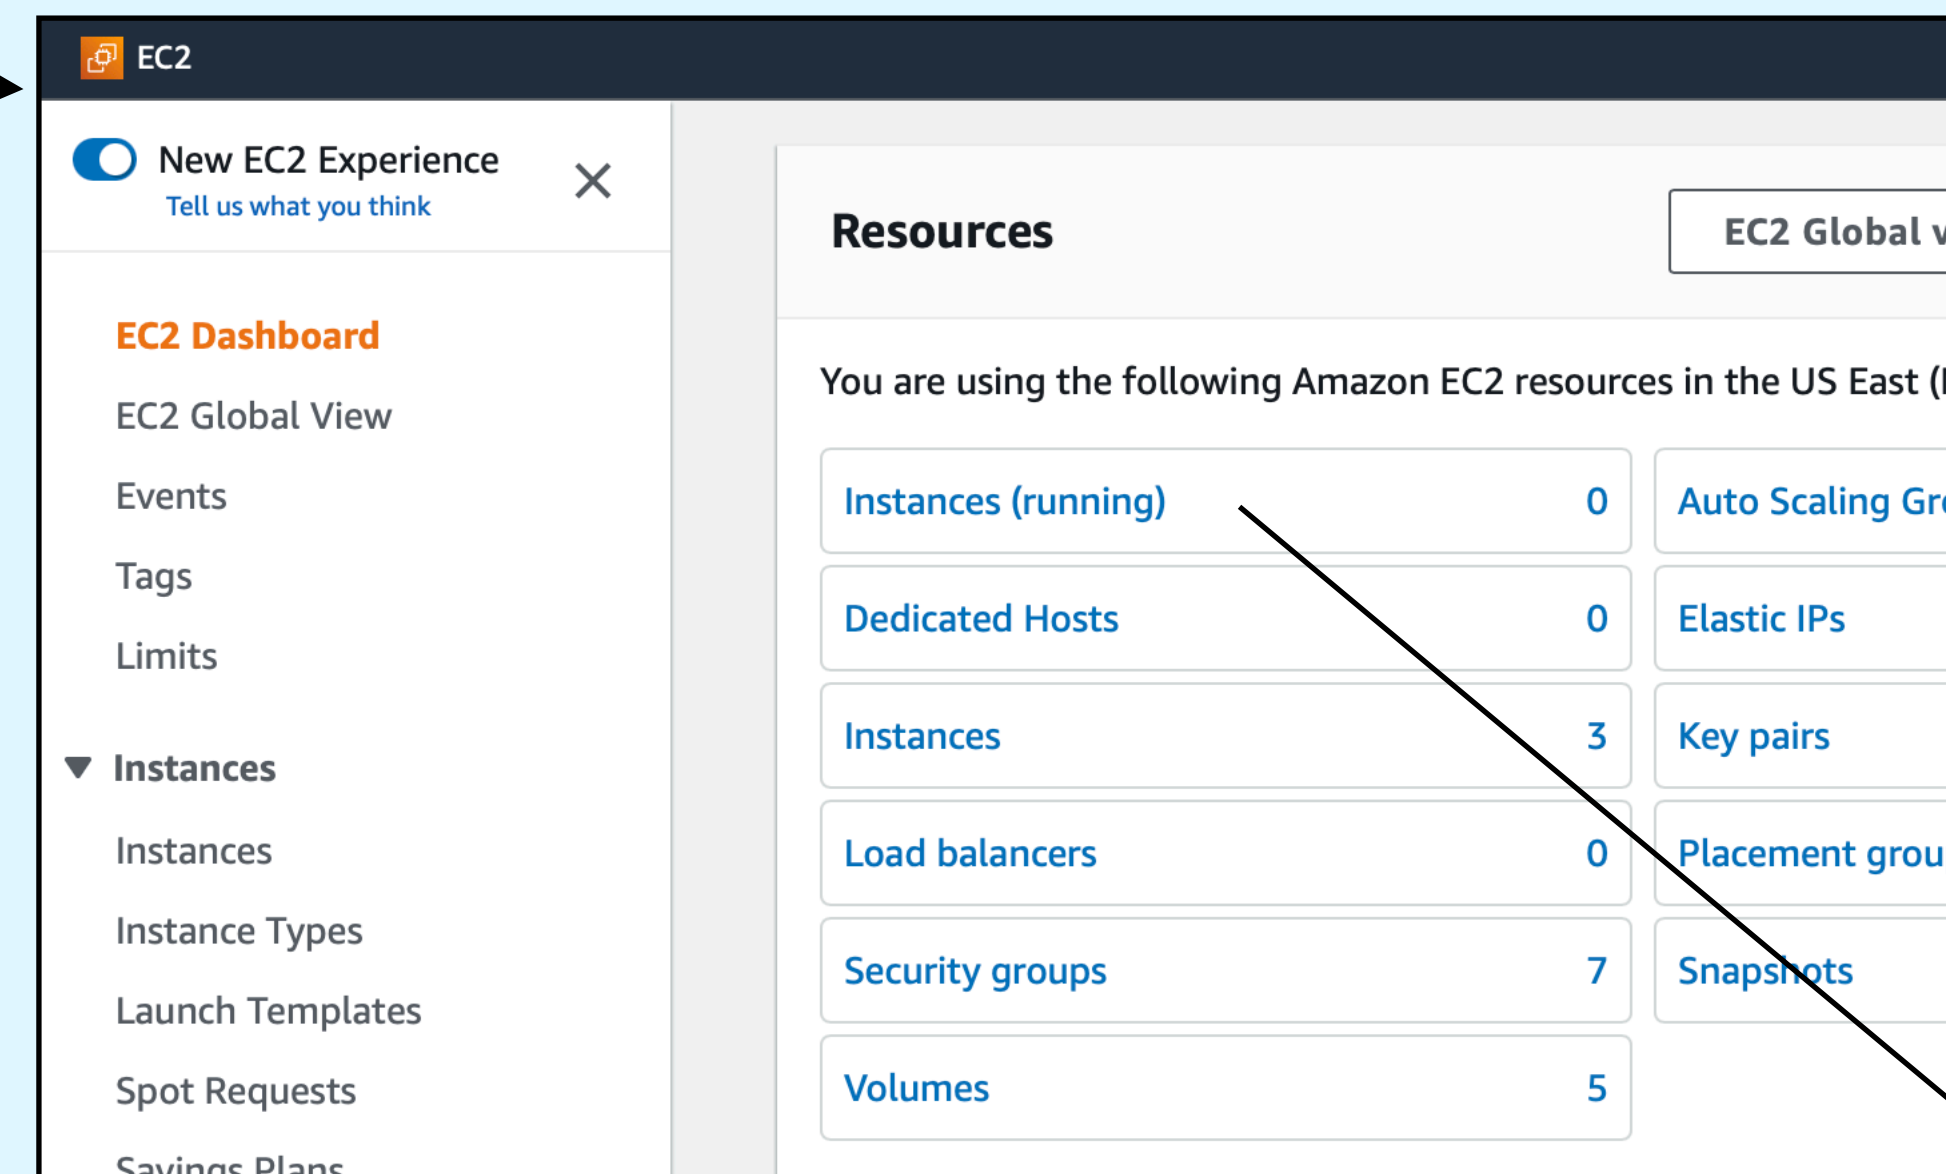
Appendix: Reference Material

- Installing the RISC-V Gnu Toolchain (gcc, as, ld, gdb, ...)
- Opening an account on Amazon AWS
- *Creating an AMI (or two) for Catamaran/ARIFIC*

Navigate to your EC2 Instances dashboard



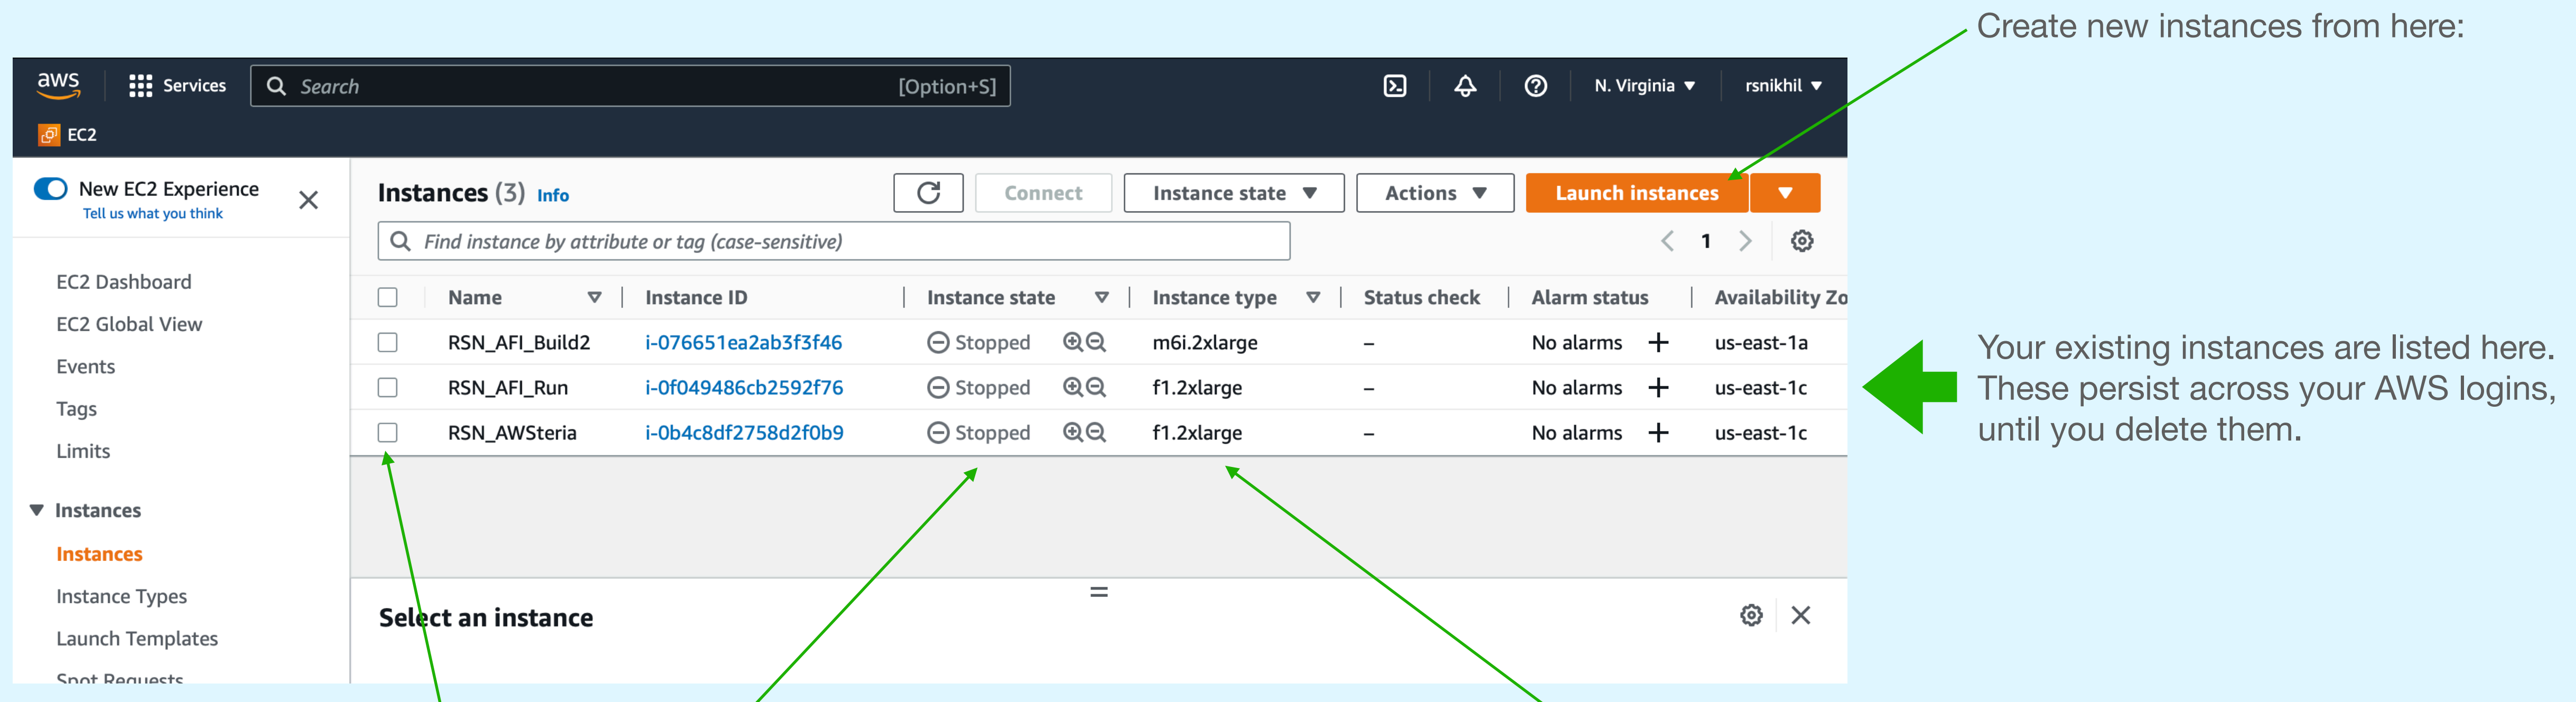
AWS offers a gazillion services; we want “EC2” (Elastic Computing), which are virtual machines in the cloud.



Then, move on to the “Instances” dashboard
 (“instances” = AWS terminology for your virtual machines)

The EC2 “instances” dashboard

Create new instances from here:



The screenshot shows the AWS Management Console for the EC2 service. The top navigation bar includes the AWS logo, a search bar, and the user's profile. The left sidebar contains navigation links for EC2 Dashboard, Global View, Events, Tags, Limits, and a dropdown for Instances. The main content area shows the 'Instances (3)' page with a search bar, a table of instances, and a 'Launch instances' button. Annotations with green arrows point to the 'Launch instances' button, the instance table, and the 'Select an instance' button.

| <input type="checkbox"/> | Name | Instance ID | Instance state | Instance type | Status check | Alarm status | Availability Zone |
|--------------------------|----------------|---------------------|----------------|---------------|--------------|--------------|-------------------|
| <input type="checkbox"/> | RSN_AFI_Build2 | i-076651ea2ab3f3f46 | Stopped | m6i.2xlarge | – | No alarms | us-east-1a |
| <input type="checkbox"/> | RSN_AFI_Run | i-0f049486cb2592f76 | Stopped | f1.2xlarge | – | No alarms | us-east-1c |
| <input type="checkbox"/> | RSN_AWSteria | i-0b4c8df2758d2f0b9 | Stopped | f1.2xlarge | – | No alarms | us-east-1c |

Your existing instances are listed here. These persist across your AWS logins, until you delete them.

Select one or more instances here

These instances are all currently “Stopped” (like “sleep” on a laptop; saves AWS charges when not being actively used)

- Selected instances can be put into the “Running” state by selecting “Start instance” in the “Instance state” drop-down menu at top (like emerging from “sleep” on a laptop).

The “Instance type” determines:

- Host CPU type (x86, ARM, ...)
- Sizing (# of virtual cores, memory, ...)
- “f1.*” types are the only ones with FPGAs attached.
- *Note:* AWS charges vary with instance type
- *Warning:* “f1” instances typically have higher AWS charges.
- *Hint:* you can do FPGA builds on an instance without an FPGA (cheaper), such as the first one shown above.

Create one or two AMIs?

The “Instance type” determines, for your AMI:

- Host CPU type (x86, ARM, ...)
- Sizing (# of virtual cores, memory, ...)
- “f1.*” types are the *only* ones with FPGAs attached.
- *Note*: AWS charges vary with instance type
- *Warning*: “f1” instances typically have higher AWS charges.

Hint (for lower AWS charges): Create two AMIs:

- A “build” AMI of type “m6i.2xlarge” for building bitfiles for Catamaran/ARIFIC
 - Use the free AWS “FPGA Developer AMI” for this, which comes preloaded with Vivado, Vivado licenses, etc. (free). The supplied OS is CentOS Linux.
- A “run” AMI of type “f1.2xlarge” for running Catamaran/ARIFIC with FPGA
 - You can choose an AMI with OS of your choice (we typically choose latest Ubuntu)

Note: when you build a bitfile on the “build” AMI, it'll reside in the cloud named by an AGFI (unique id), which you can access from the “run” AMI

“Create” your virtual machine (“instance”)

When you select “Launch Instances” to create a new instance, it will take you through a series of steps (“Launch Instances Wizard”).

Here are some notes on the various steps:

- “Name and tags”: choose a name for your instance (e.g., on our slides: “RSN_AFI_Build2”, “RSN_AFI_run”, “RSN_AWSteria”)
- Application and OS Images (Amazon Machine Image)
 - Choose one of the AMIs from the 1000s available in the search box
 - For Catamaran/ARIFIC bitfile building, we recommend the latest AWS “FPGA Developers AMI”
 - For Catamaran/ARIFIC running on FPGA, we recommend an AMI with the latest Ubuntu (22.04 LTS at time of writing), x86 (not ARM!), 64-bit
- “Select and Instance Type”
 - The menu shows 100s of possibilities, x86 and ARM, etc.
 - For Catamaran/ARIFIC bitfile building, we recommend “m6i.2xlarge” (e.g., see “Instance type” on previous slide)
 - For Catamaran/ARIFIC running on FPGA, this *must* be an “f1” instance; we recommend “f1.2xlarge” (e.g., see “Instance type” on our slides)
- “Create your Key Pair”
 - When one connects to an AMI (with “ssh”), for security reasons Amazon does not support password-based login authentication, only crypto-key based authentication.
 - This menu item creates a key pair for you, a public key and private key.
 - The prompts will ask you to copy the private key, and save it on your laptop/desktop in a file `~/.ssh/MyPrivateKey.pem` and set its protection to 0x400
 - Note: once you’ve created a key pair, you can reuse it for new instances; it will show you your existing key pairs and you can select one, instead of creating a new one.
- “Network Settings”: Leave unchanged. In particular, “Auto-assign public IP” = “Enable” and “Allow SSH traffic from Anywhere”
- “Configure Storage”: Leave unchanged, e.g., 8 GiB (these can be updated later, if needed)
- “Advanced Details”: Leave unchanged
- Finally, select “Launch Instance” on the “Summary” panel which floats on the right during the wizard.
 - (but see caveat on next slide)

“Create” your virtual machine (“instance”)

Caveat: After selecting “Launch Instance” at the end of the AMI-creation wizard, you may encounter an error message like this:

You have requested more vCPU capacity than your current vCPU limit of 0 allows for the instance bucket that the specified instance type belongs to. Please visit <http://aws.amazon.com/contact-us/ec2-request> to request an adjustment to this limit.

If you get this message, please visit the contact-us link shown, click on “Support” and lodge a request to increase your “vCPU quota” to 8, which is required for the f1.2xlarge instance type. Try to include a few sentences explaining your need for an f1.2xlarge instance type.

You will get an almost instance response saying:

This specific limit increase request requires further internal review before approval ...

These requests are processed by humans, so it may take a few hours before you receive an email indicating that your request has been approved, after which you can proceed.

END

GitHub repo for this tutorial: <https://github.com/rsnikhil/Tutorial> at HPCA-29